Lambda Architecture

Santosh Hari

Santosh Hari

Azure Consultant @ Nebbia Tech

Azure MVP

President, Orlando .NET UG

Organizer, Orlando Codecamp

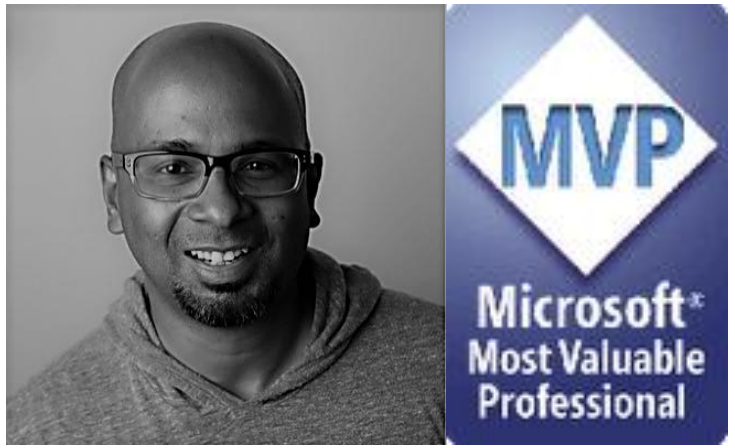

santoshhari.wordpress.com

santosh.hari@newsignature.com

@_s_hari

/in/santoshhari

Agenda

- Intro to Lambda Architecture (LA)
- Intro to Cosmos DB
- Discussion on building blocks
  - Materialized View
  - Event Sourcing
  - CQRS
- Big picture view
- Discussion on how Cosmos DB simplifies LA
- Criticisms of LA and short discussion on Kappa Architecture
- Heavy on concepts, no code or demo

Intro to Lambda Architecture

Modern data system scale horizontal

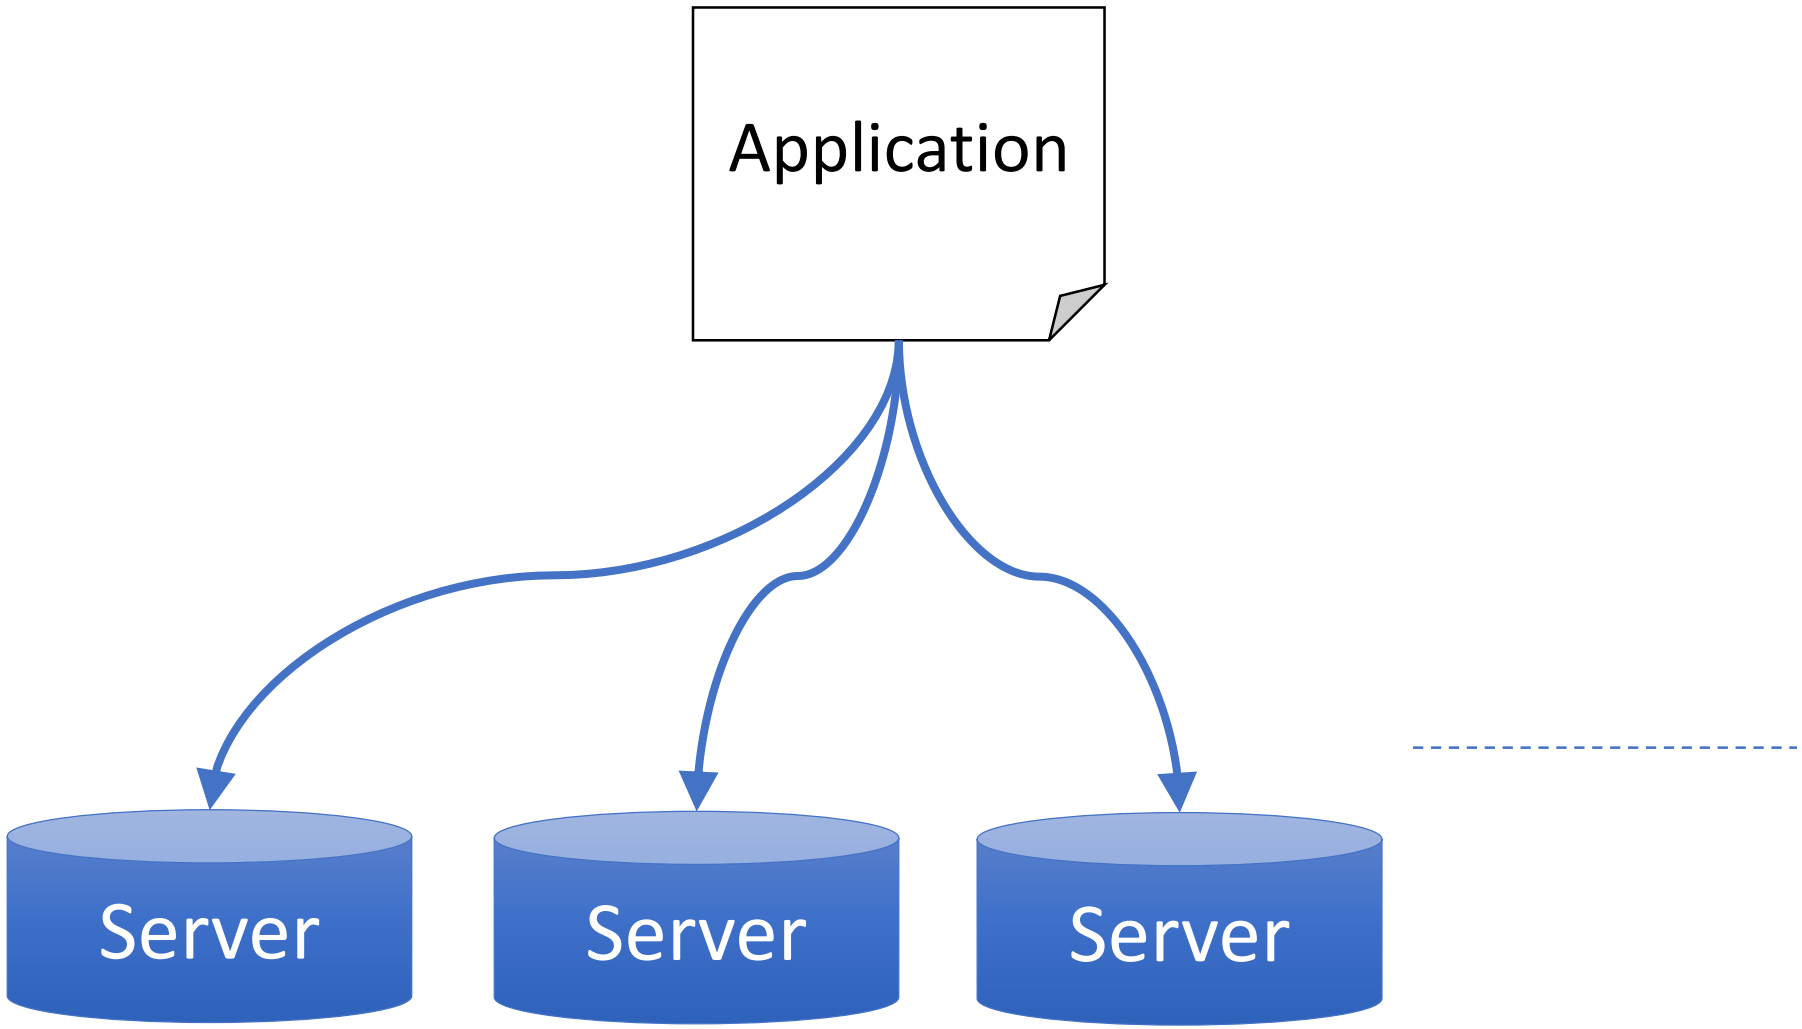

CAP Theorem

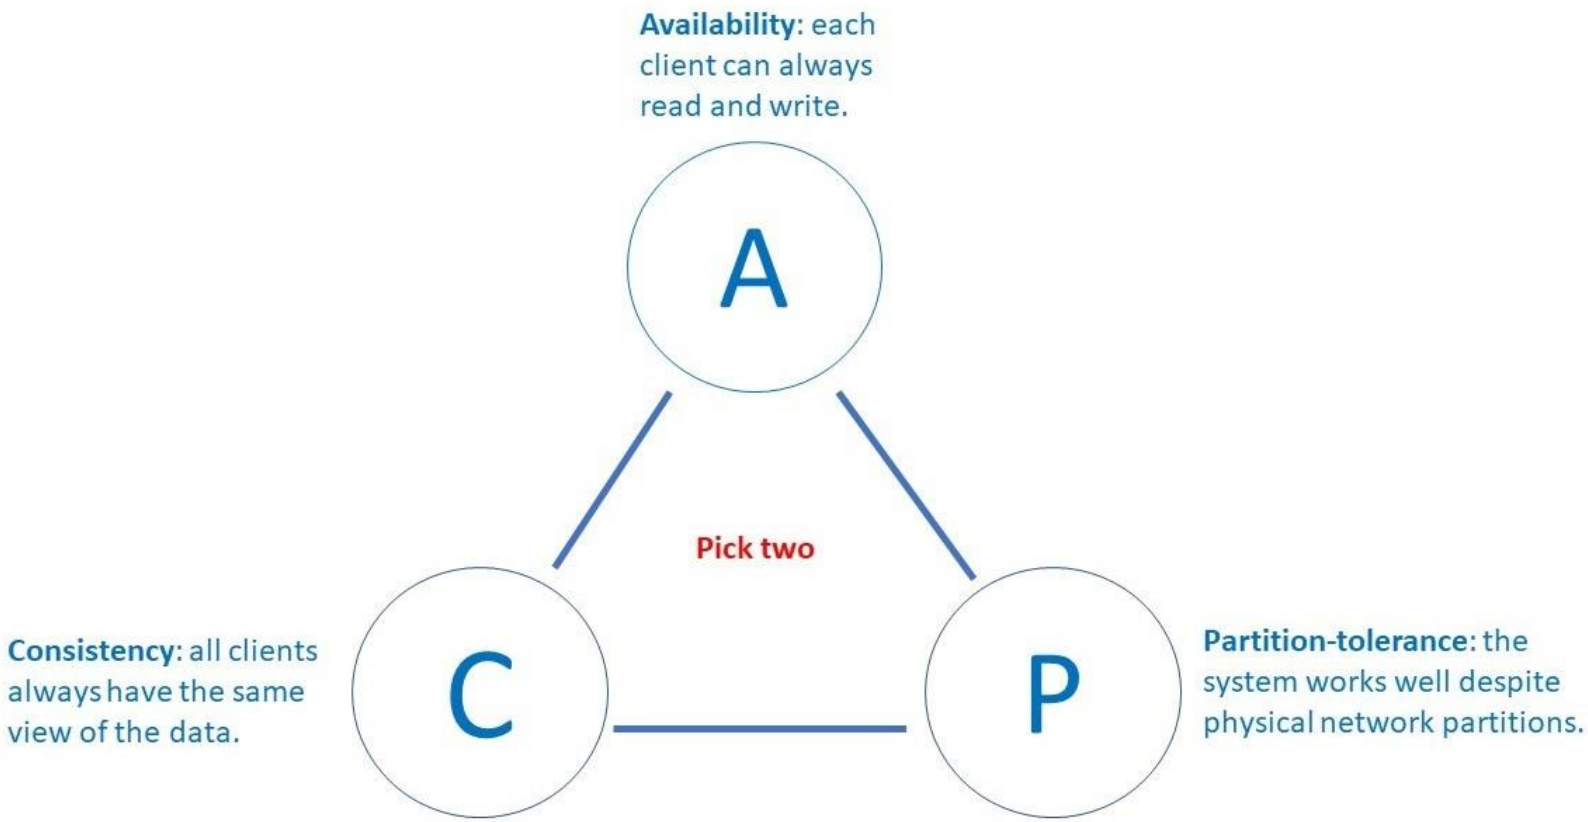
Availability: each client can always read and write.
A
Pick two
Consistency: all clients always have the same view of the data.
C
P
Partition-tolerance: the system works well despite physical network partitions.

Challenge of data processing in distributed systems

- High latency
  - Unable to read latest writes
  - Keep data in it’s original form

- Low latency
  - Issues querying large amounts of historical data
  - Fault tolerance
  - Not very scalable

What is Lambda Architecture?

Term coined by Nathan Marz in 2012

Distributed data processing architecture

Generic, scalable

Robust, fault tolerant against system and human failure

Enables low latency reads and updates

Scales horizontally

Lambda Architecture Overview

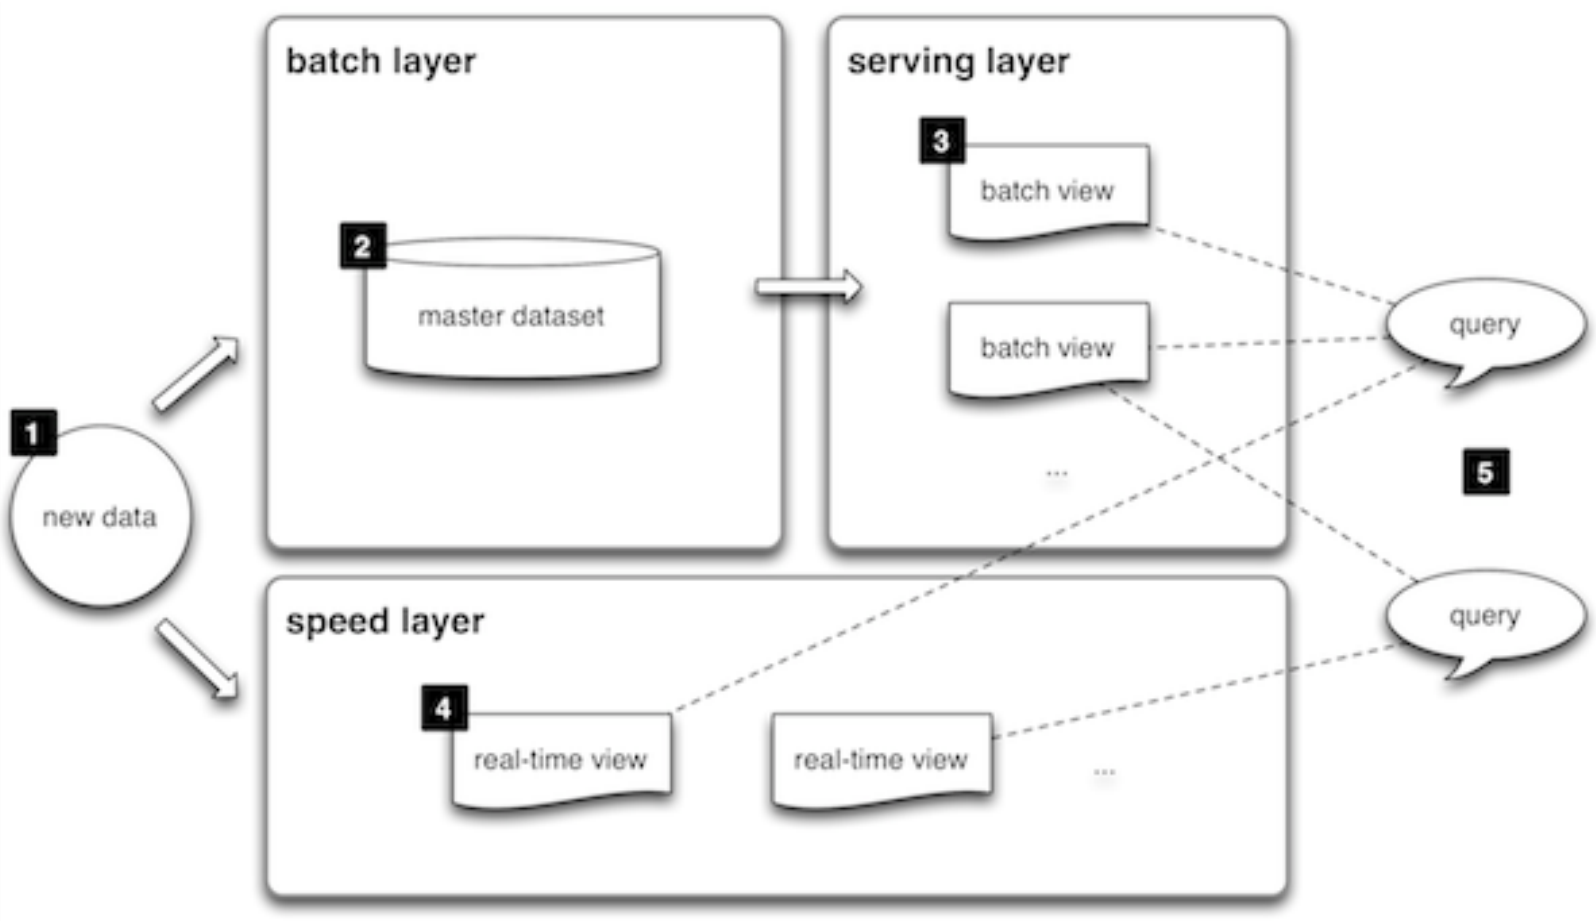

Use Case Scenarios

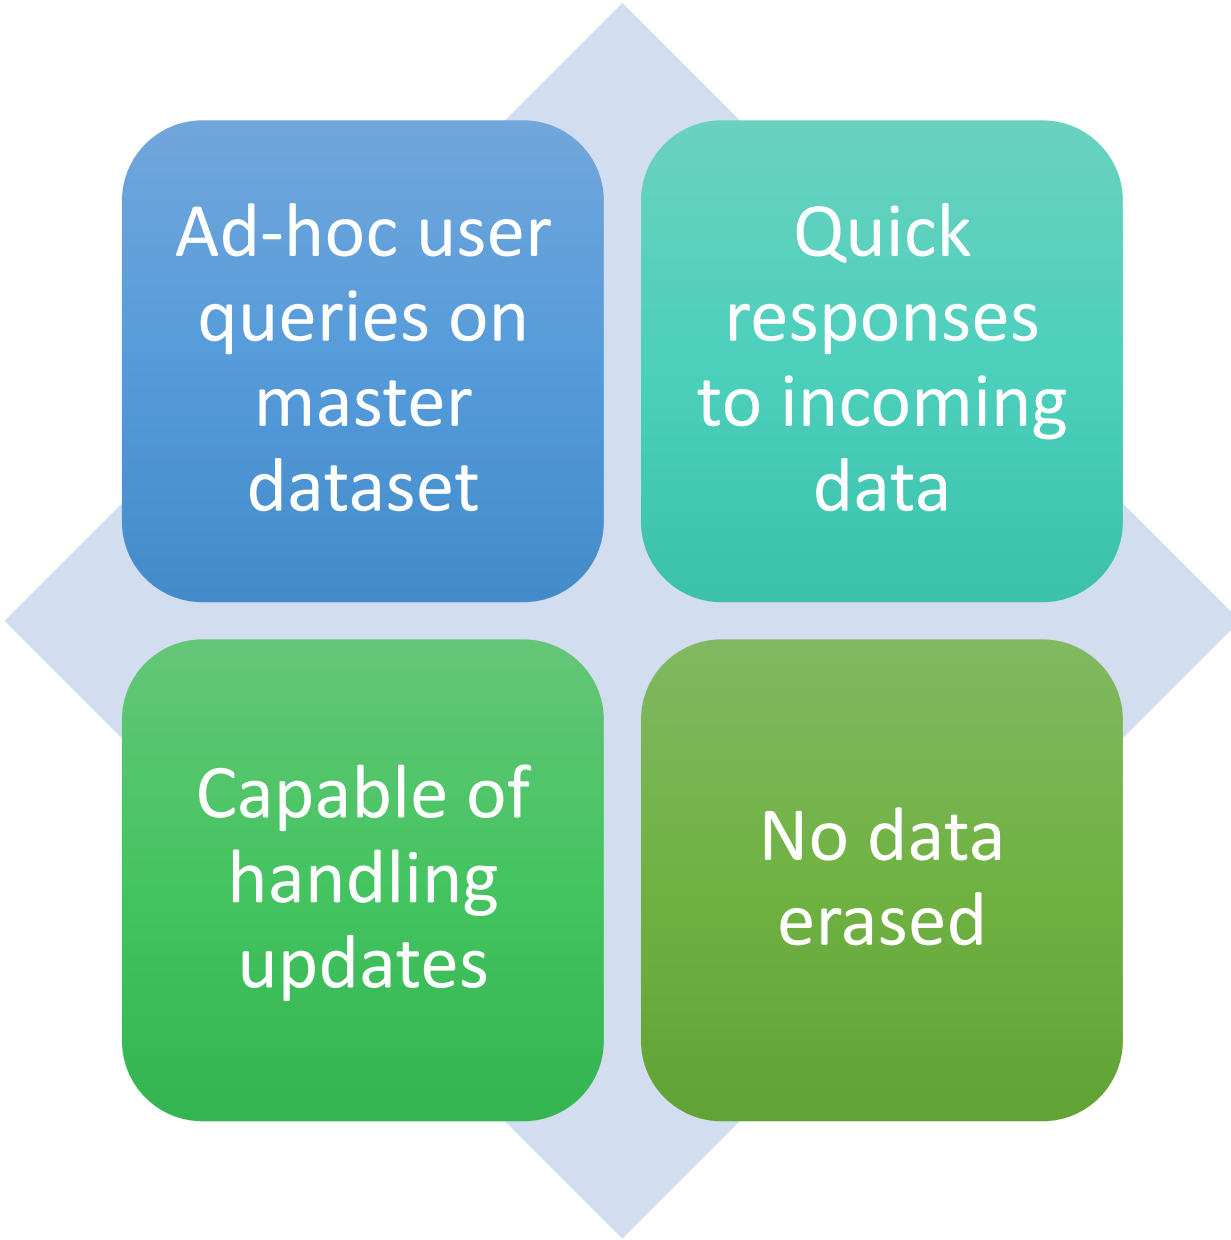

Pros and Cons

CONS

- Complex
- Re-processes every batch cycle
- Difficult to migrate / reorganize

PROS

- History
- Less errors
- Speed and reliability.
- Fault tolerant
- Scalable


Cosmos DB

Enter Cosmos DB

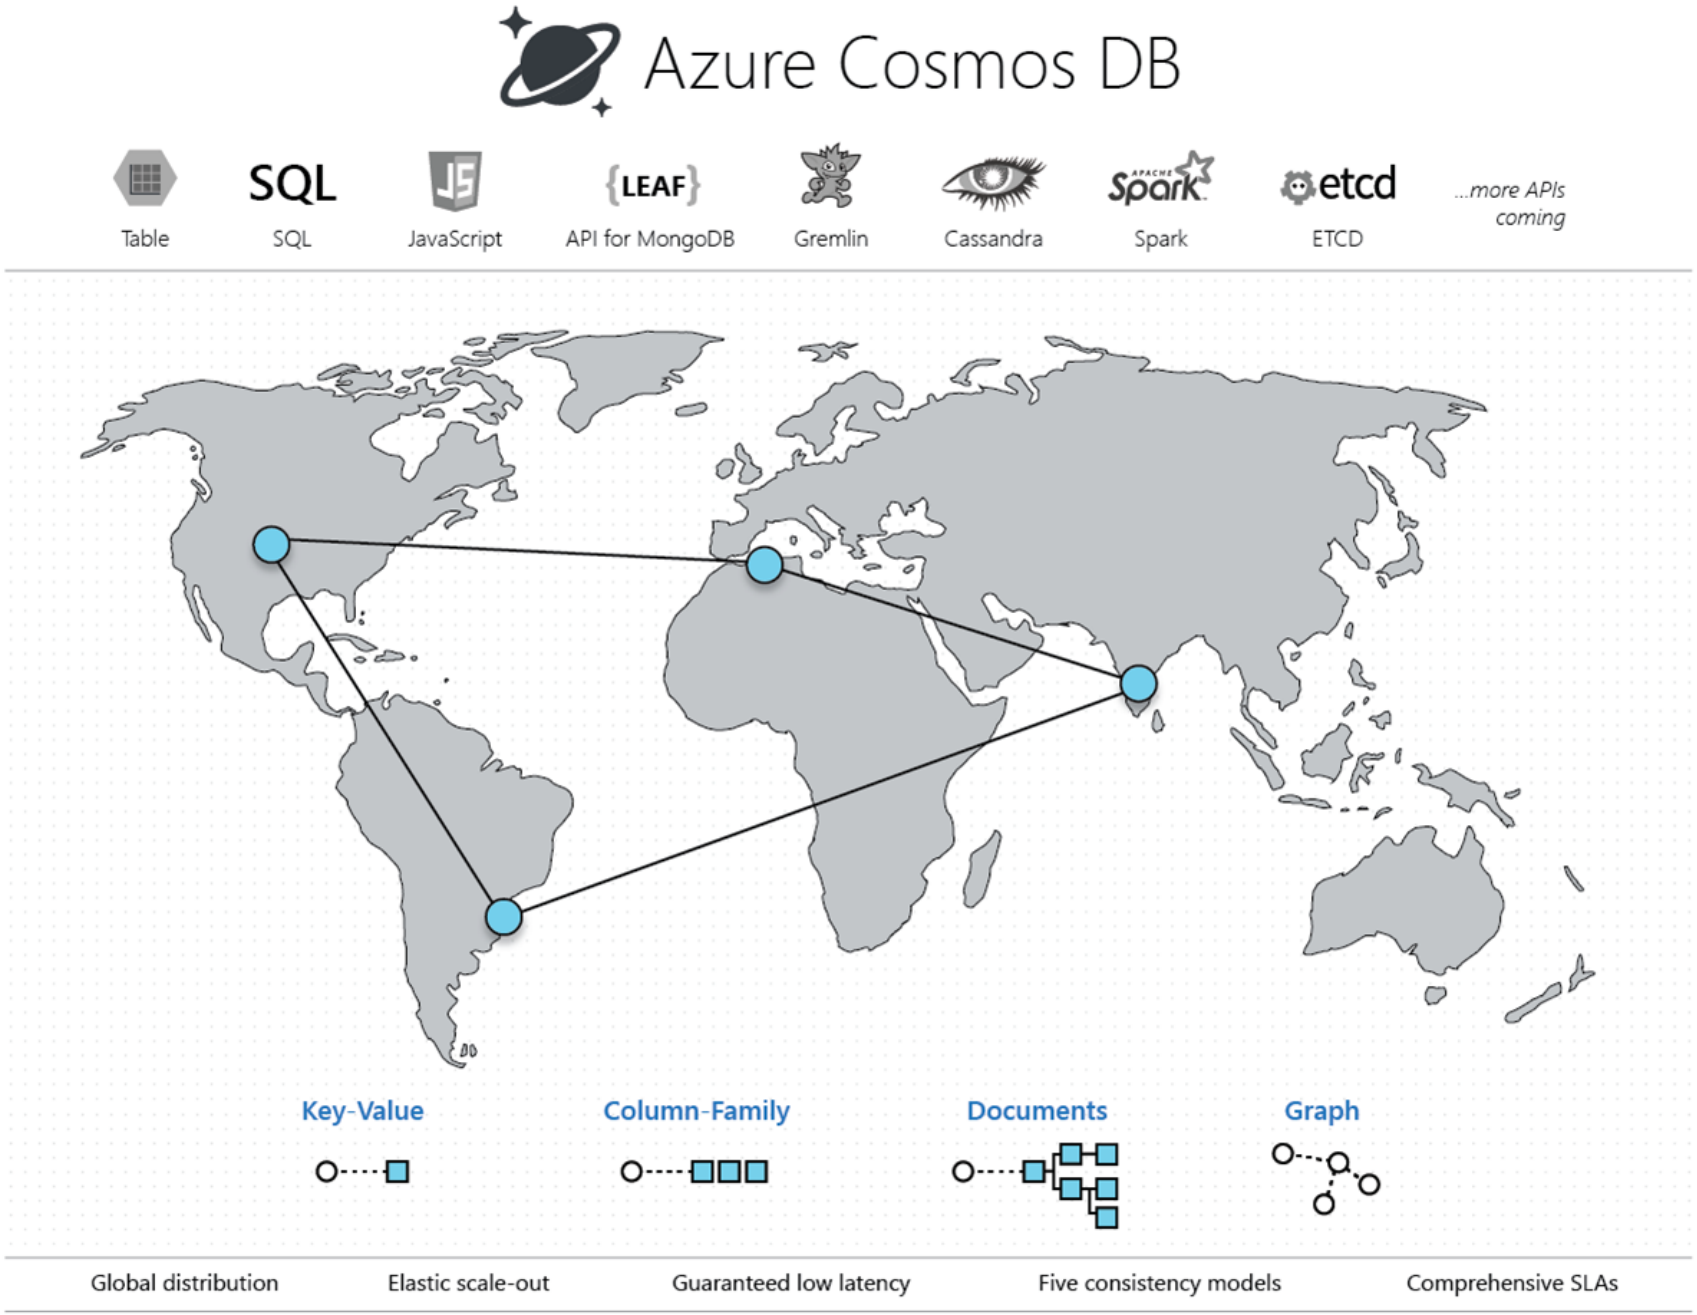

Cosmos DB Change Feed

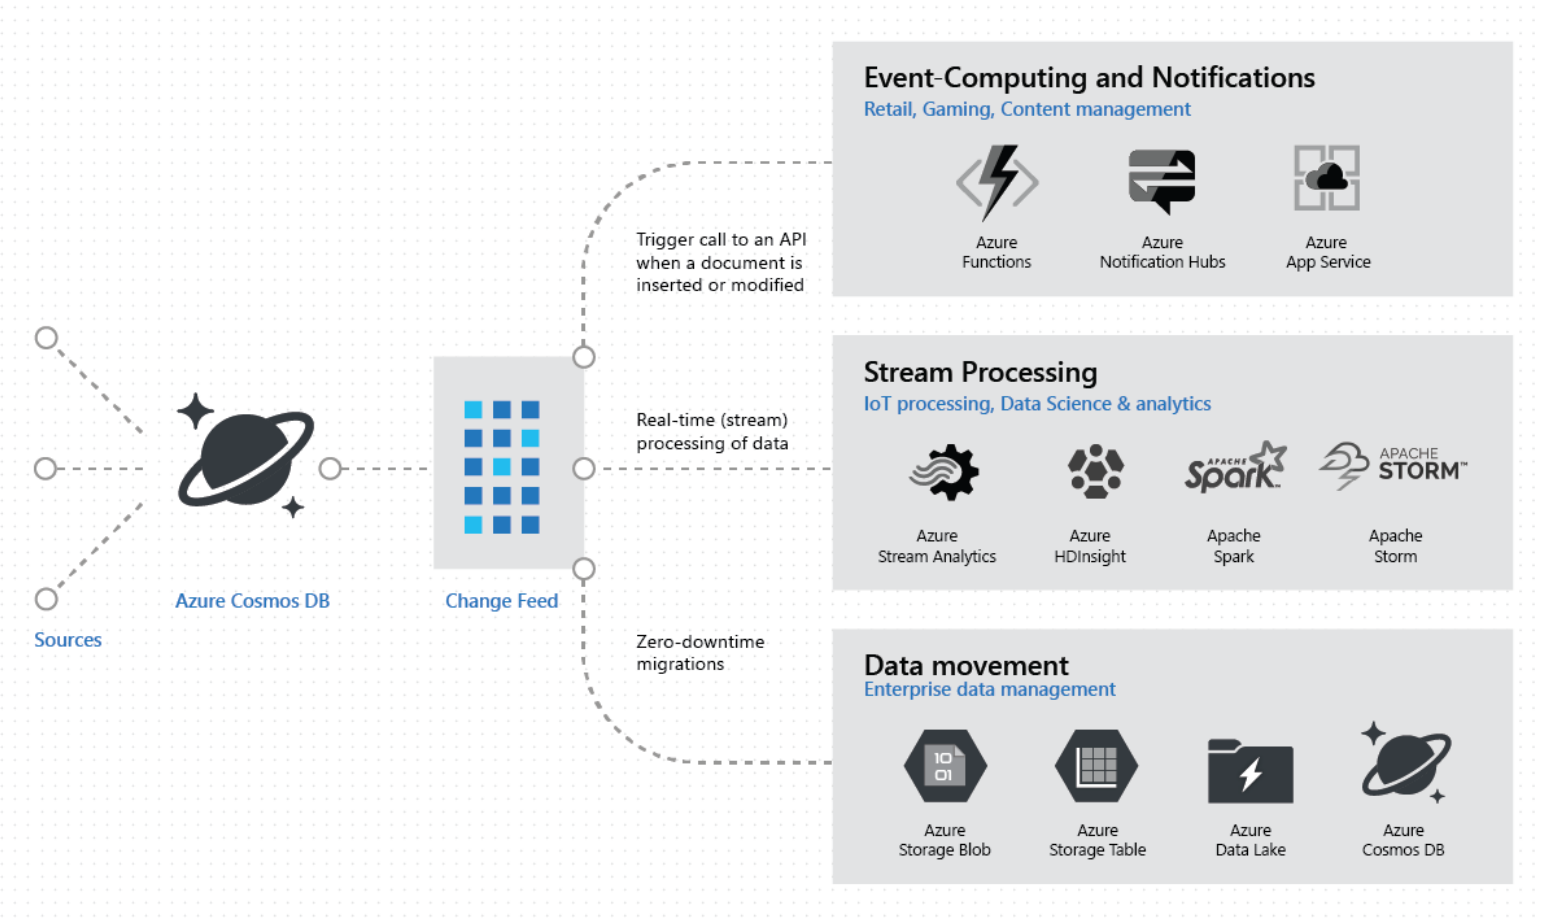

Relevant Patterns

Patterns to know

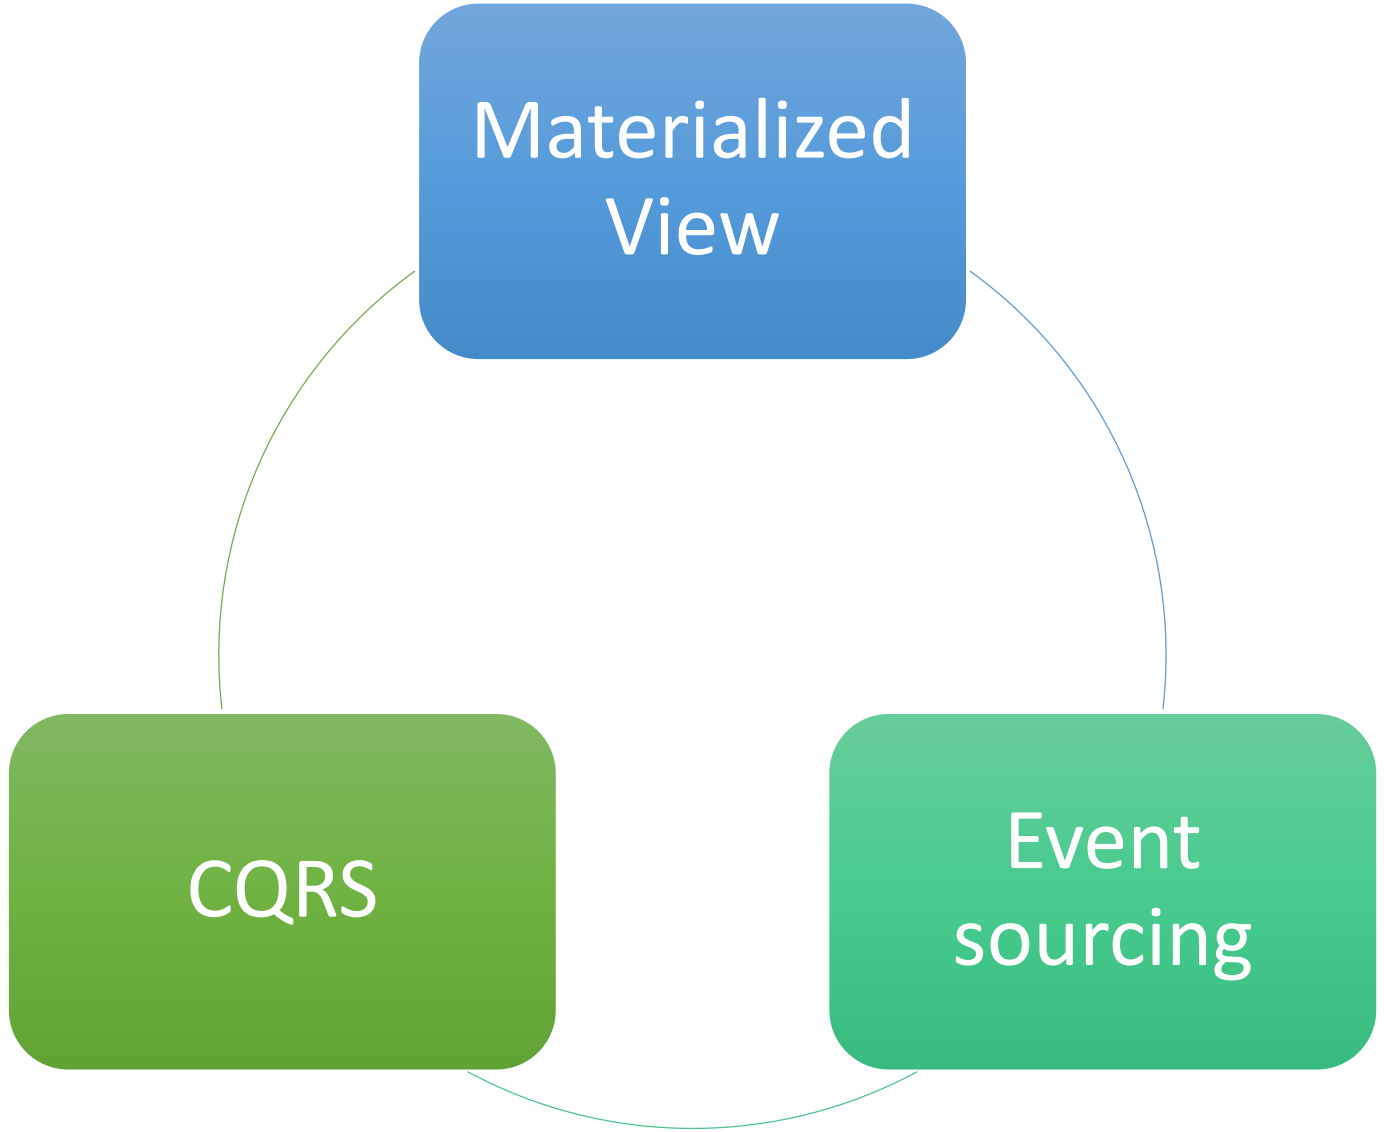

Materialized View

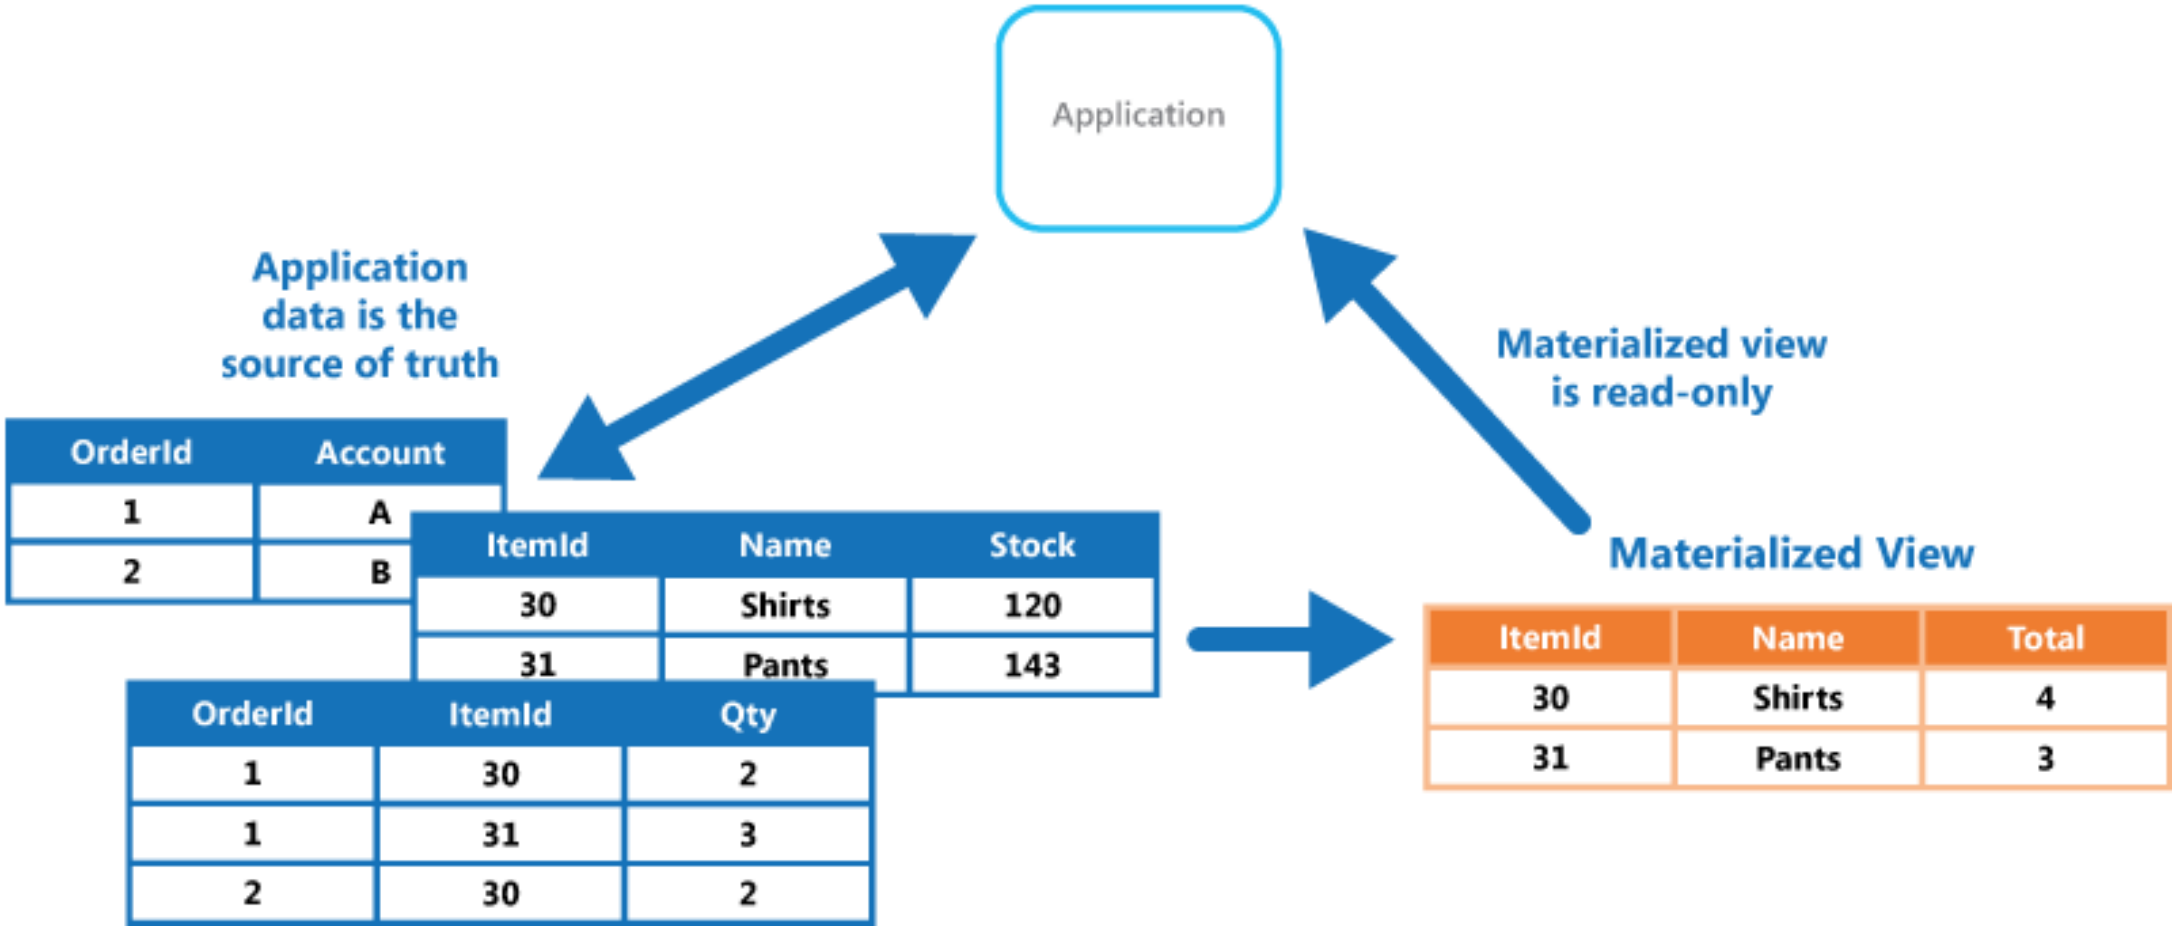

What problems does Materialized View solve

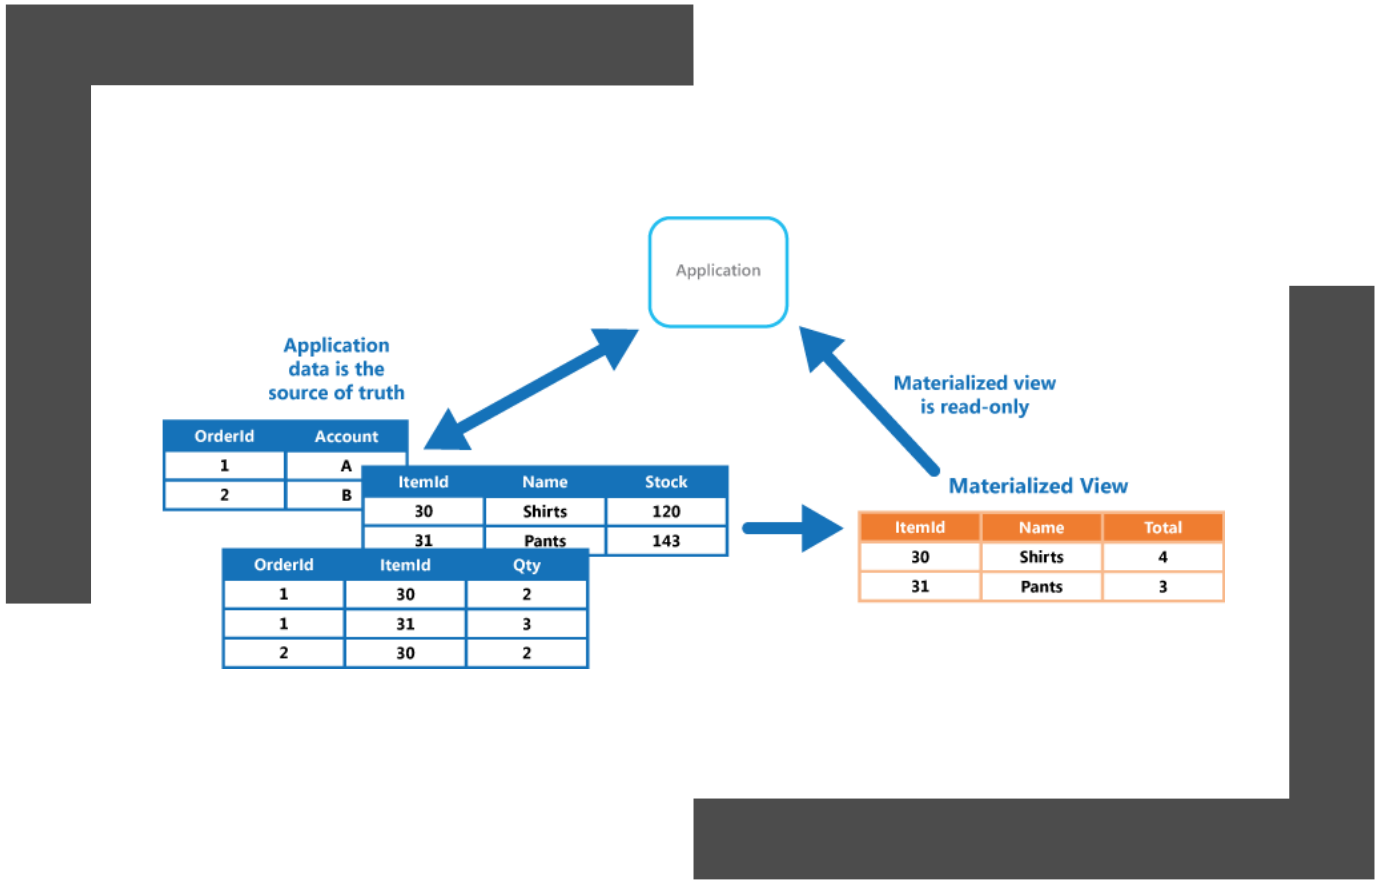

- Data optimized for storing or reading
- Negative effect on queries

How does Materialized View solve these problems

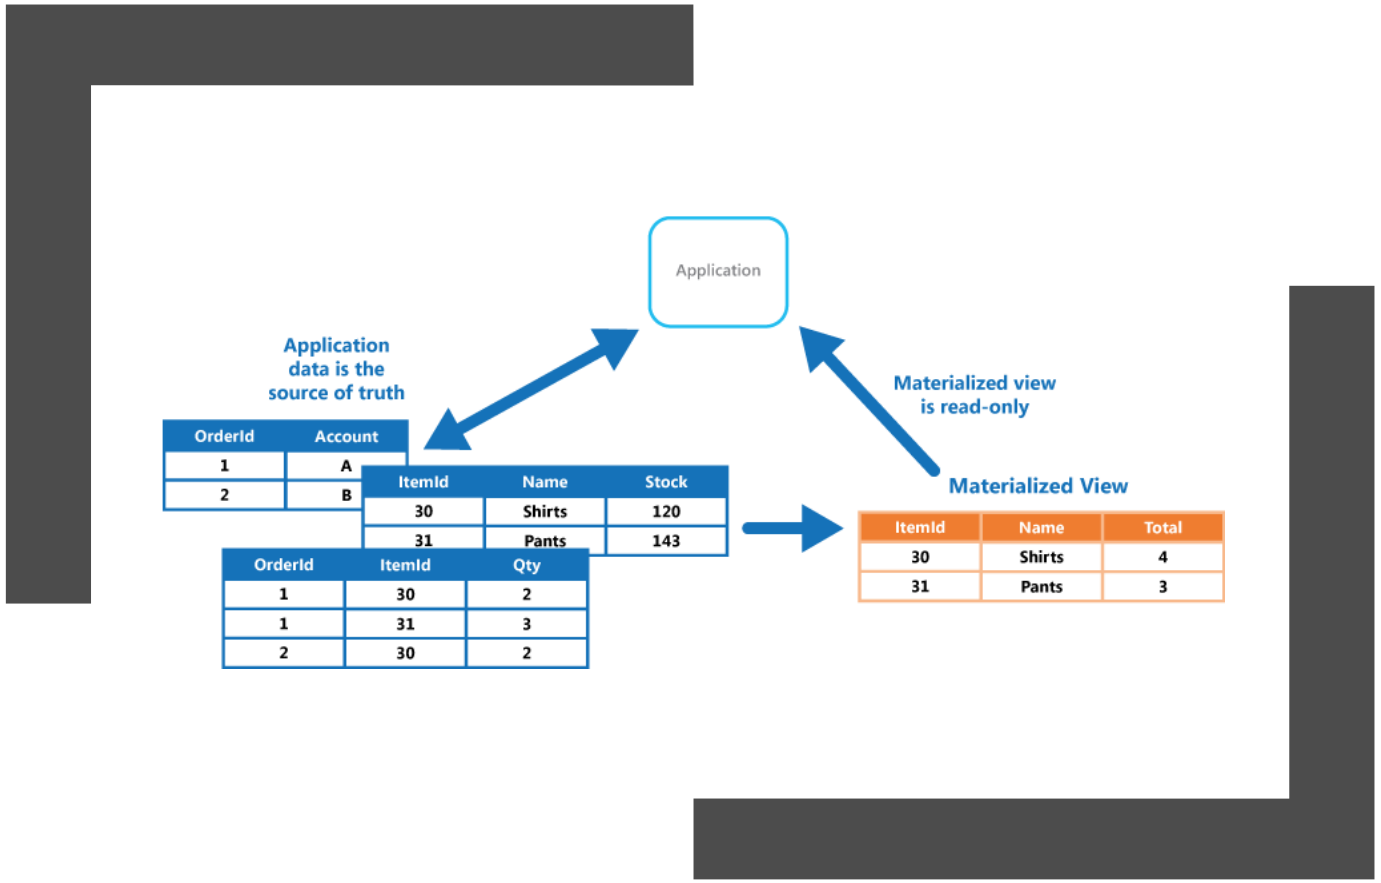

- Contain only read data
- Subset
- Disposable
- Easy re-generation
- Cache

Materialized View considerations

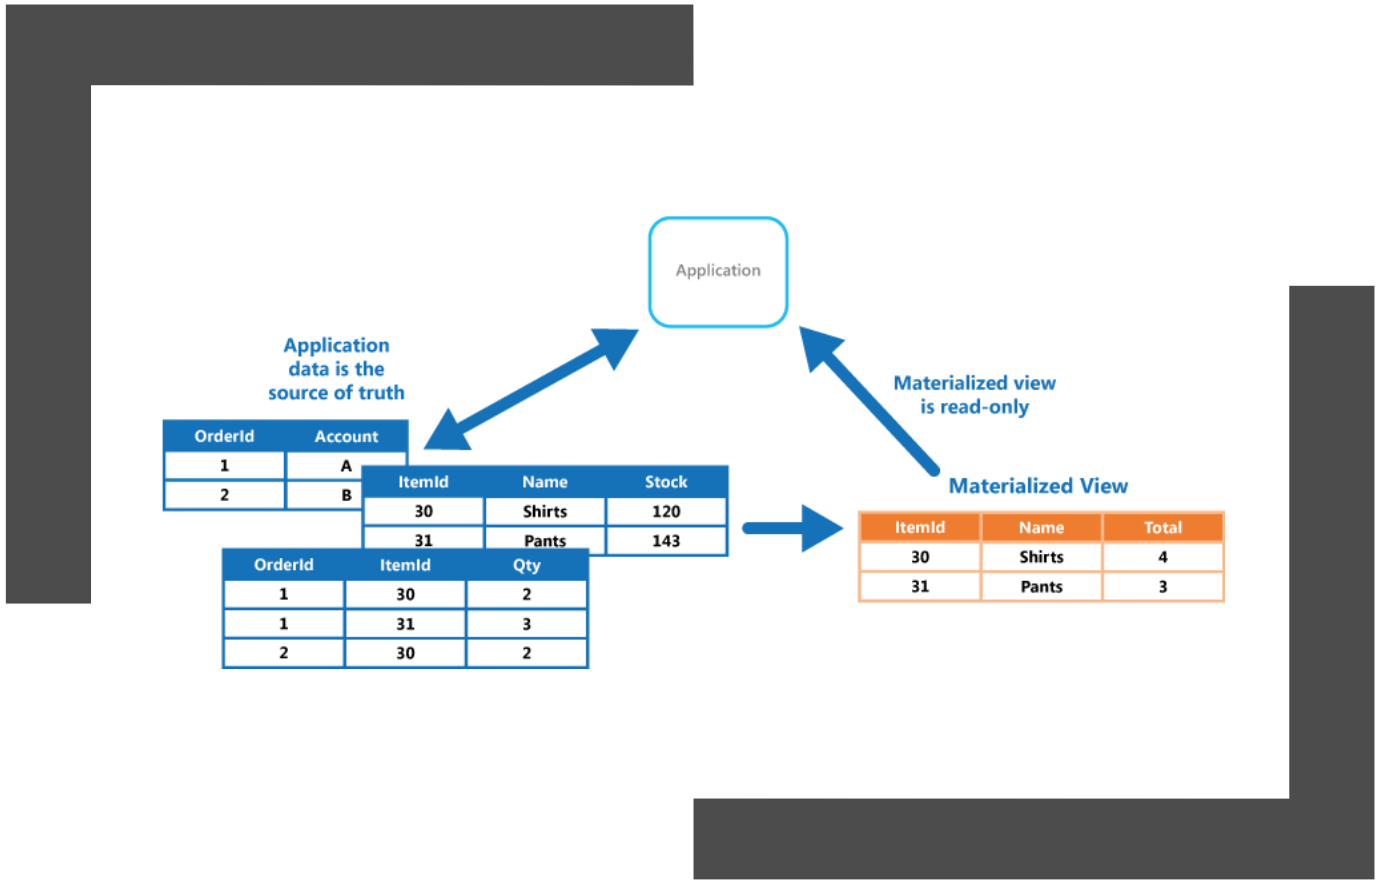

- Handling updates
- **Eventual consistency**
- Storage
- Index and partitioning

Materialized View use cases - good

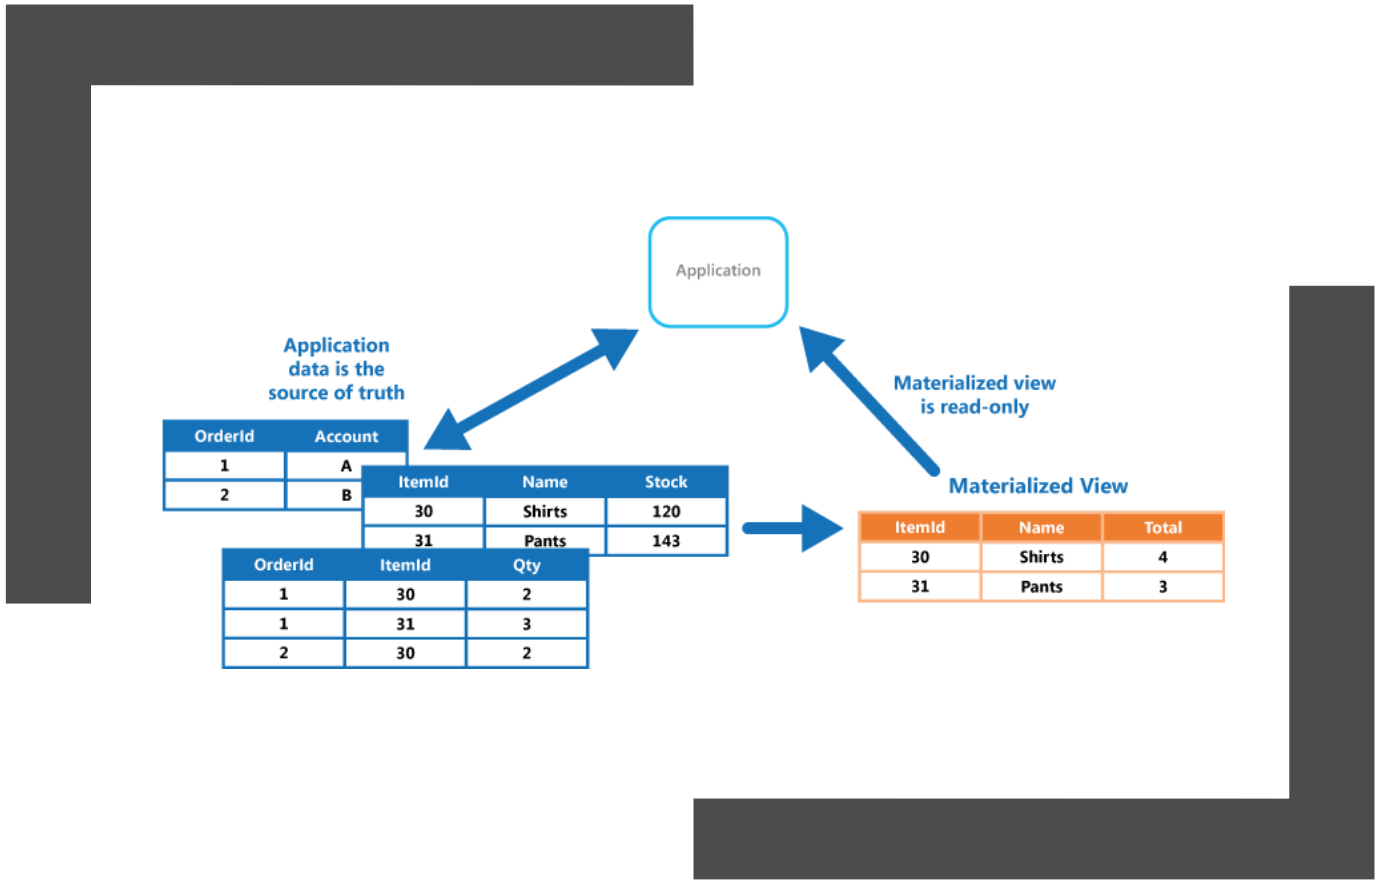

- Handles difficult to query data
- Perf
- Local cache
- Separate query from source
- Bridging data sources

Materialized View use cases - bad

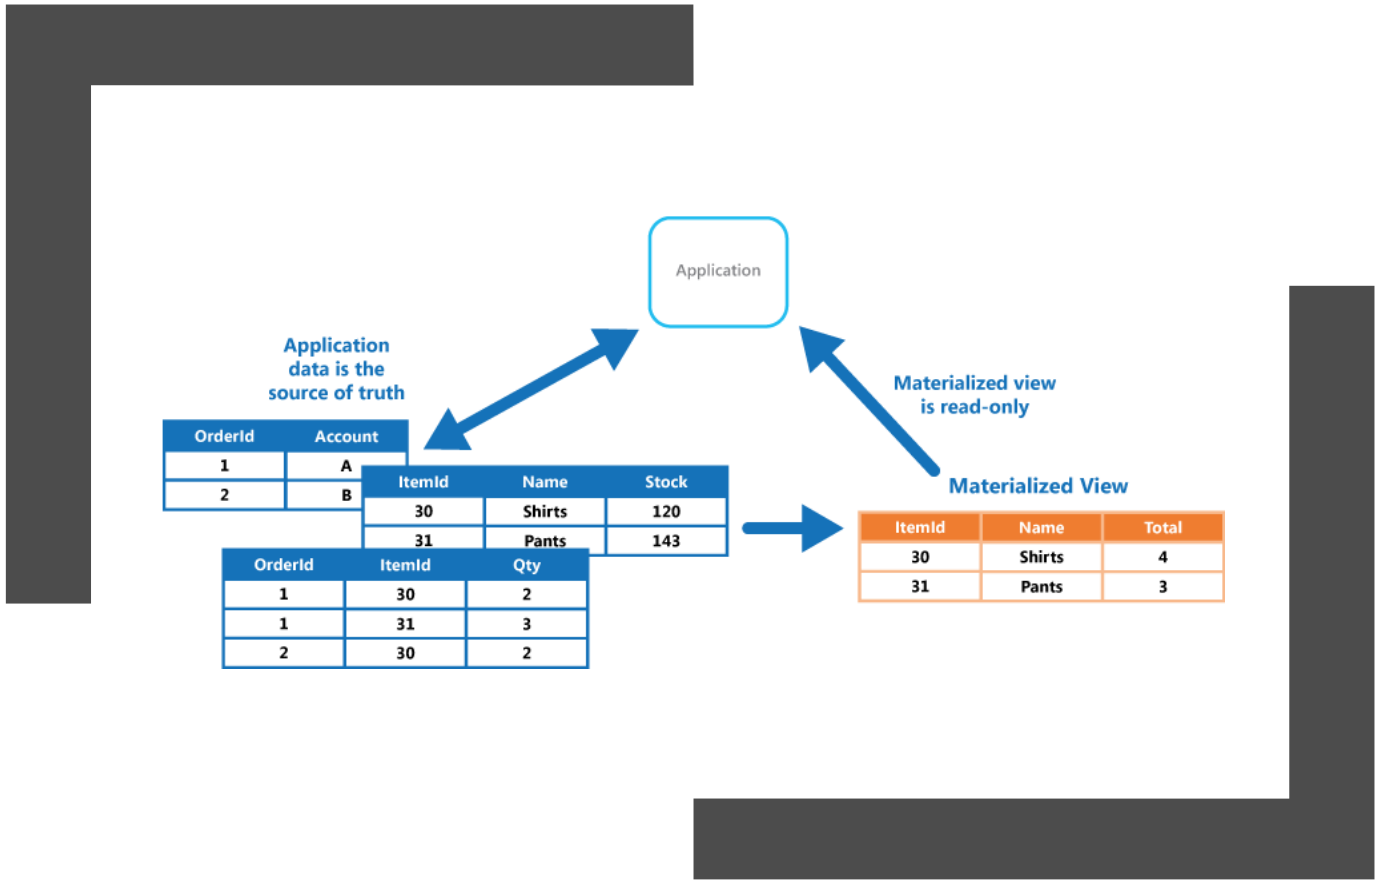

- Simple to query data
- Rapidly changing data
- Consistency needed

Materialized View with Cosmos D

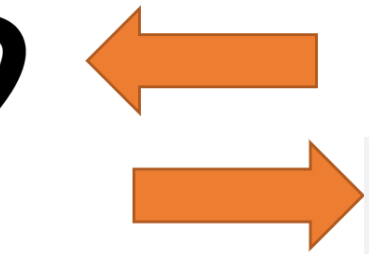

```
{
    “deviceId”:”030”,
    “currentValue”: 155.4189
}
```

```
{
    "id": "global",
    "deviceId": "global",
    "type": "global",
    "deviceSummary": {
        "035": 104.3423159533843,
        "016": 129.1018793494915,
        ...
        "023": 177.62450146378228,
        "033": 178.97744880941576
    },
}
```

```
{
    "id": "030",
    "aggregationSum": 3519.8782286699293,
    "lastValue": 155.41897977488998,
    "type": "device",
    "deviceId": "030",
    "lastUpdate": "2019-03-22T19:50:17Z"
}
```

```
foreach(var d in input)
{
    var device = Device.FromDocument(d);

    var tasks = new List<Task>();

    tasks.Add(p.UpdateDeviceMaterializedView(device));
    tasks.Add(p.UpdateGlobalMaterializedView(device));

    await Task.WhenAll(tasks);
}
```

https://medium.com/@mauridb/real-time-materialized-views-with-cosmos-db-90ece8847b50

Event Sourcing

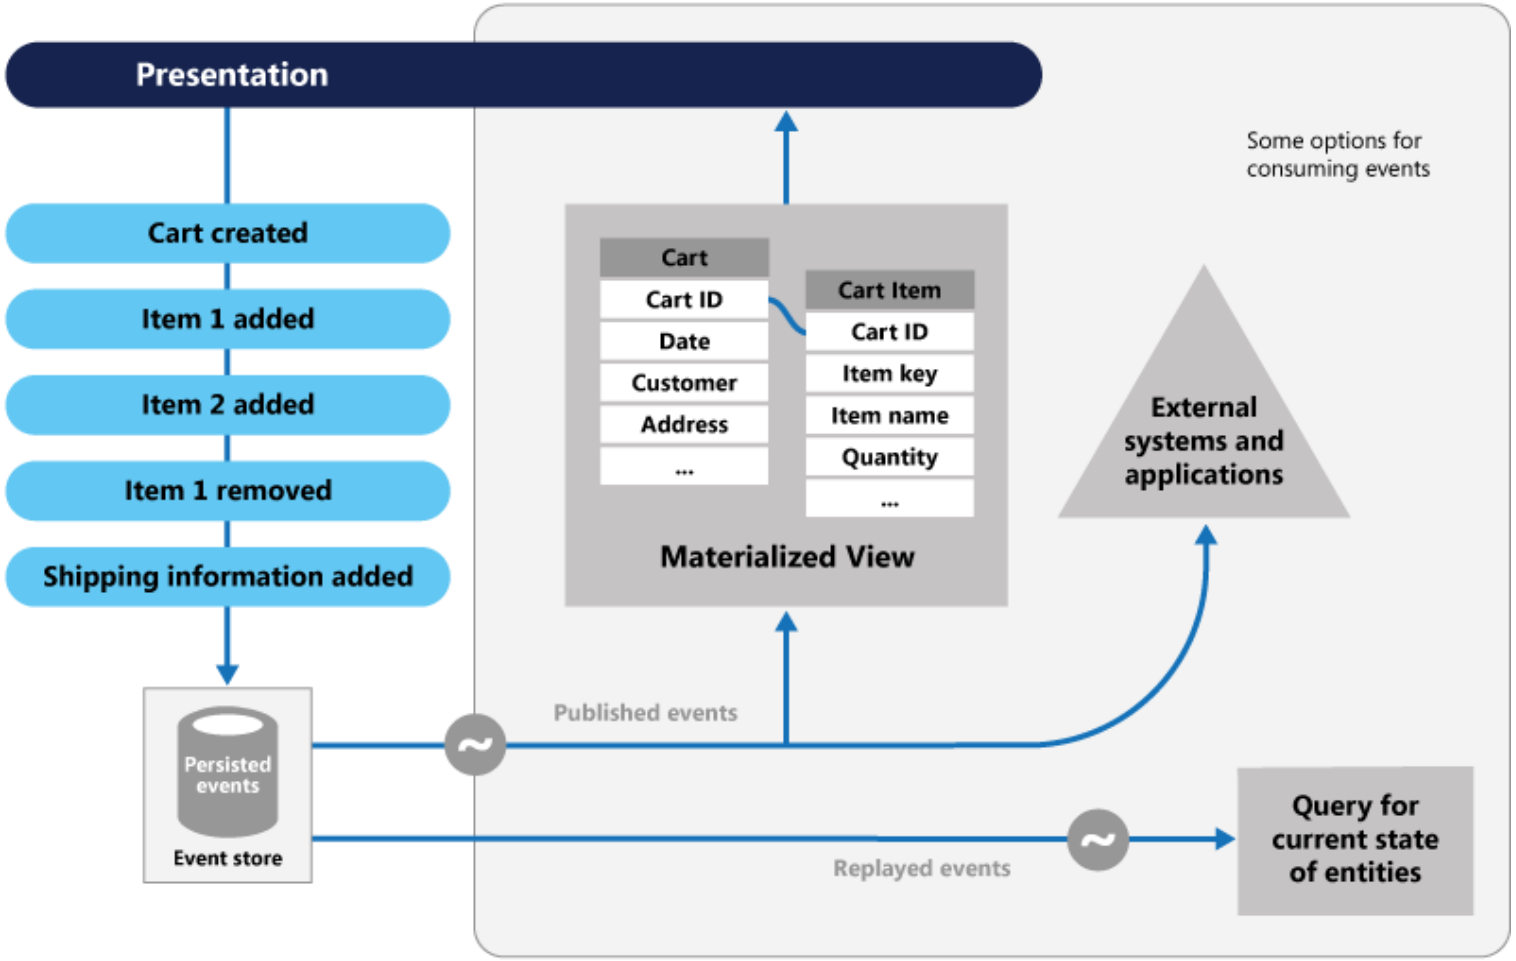

What problems does Event Sourcing solve?

- Current state only
- Locking
- Data loss
- Audit??

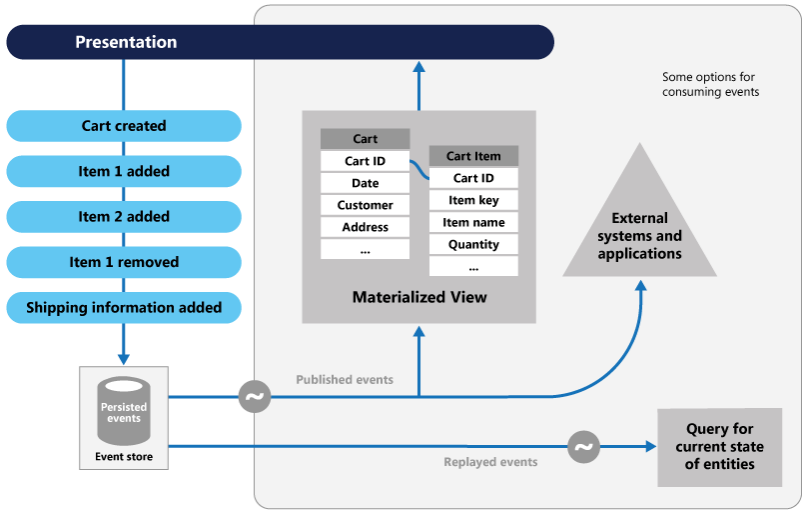

How Event Sourcing solves these problems

- Data-as-events-sequence
- Append-only store
- Store is authoritative
- Store publishes events
- Materialized View
- Replay events

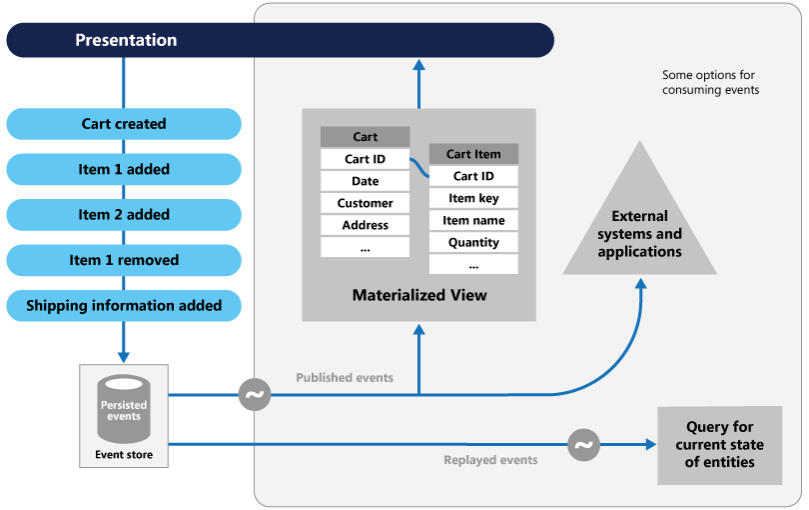

Event Sourcing advantages

- Immutable events
- Simple objects
- Domain experts friendly
- No direct DB updates
- Audit trail
- Decouple events from tasks

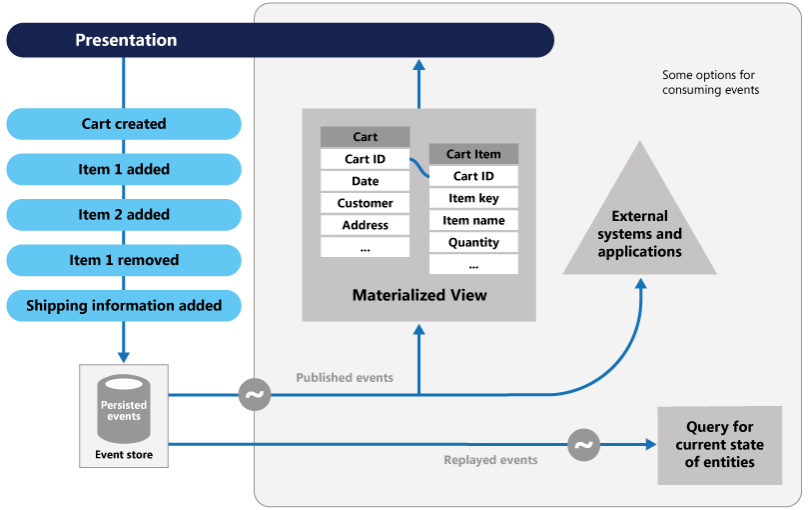

Event Sourcing considerations

- **Eventual Consistency**
- No out-of-band updates to event store
- Event format vital
- Order of events vital
- Current State = sum all
- Use snapshots
- Reduces data conflicts
- Consumers idempotent

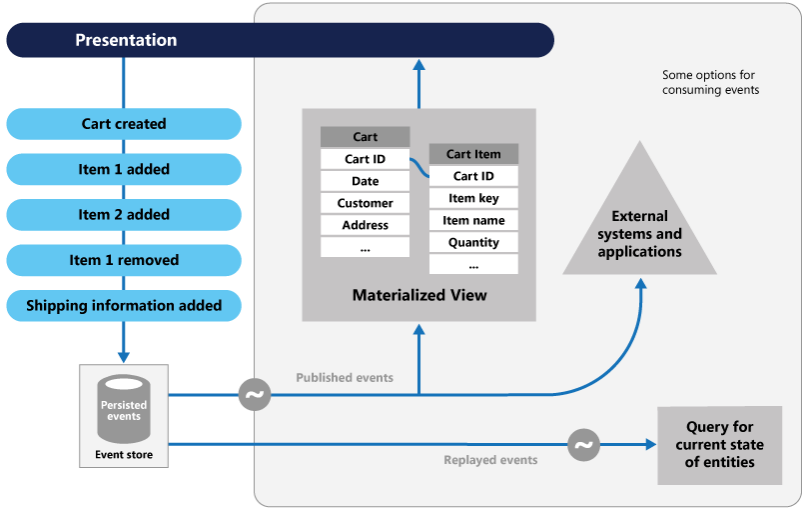

Event Sourcing use cases - good

- Capture intent, purpose, reason
- Avoiding conflicts
- Restore / rollback
- Decouple I/P O/P
- Great with CQRS

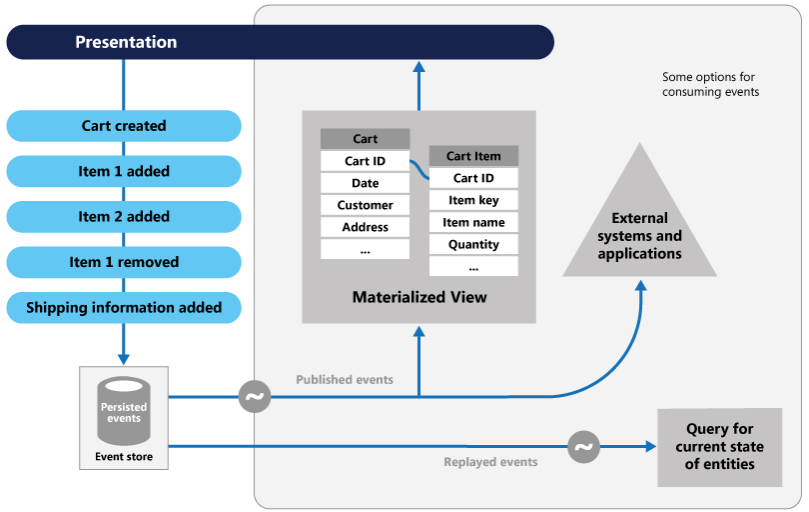

Event Sourcing use cases - bad

- Simple domain
- CRUD
- Consistent systems
- Real-time systems
- Non-audit non-historical systems
- Low conflict systems

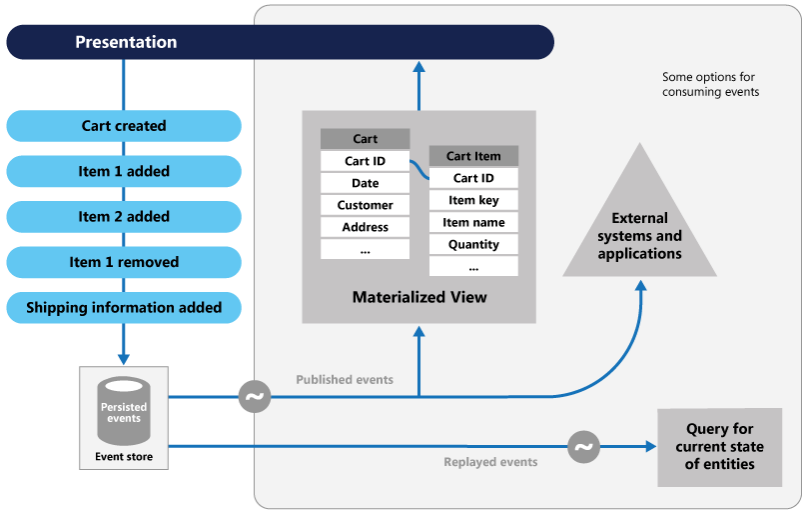

Event Sourcing with Cosmos DB

{“MessageId”:1547632386819}

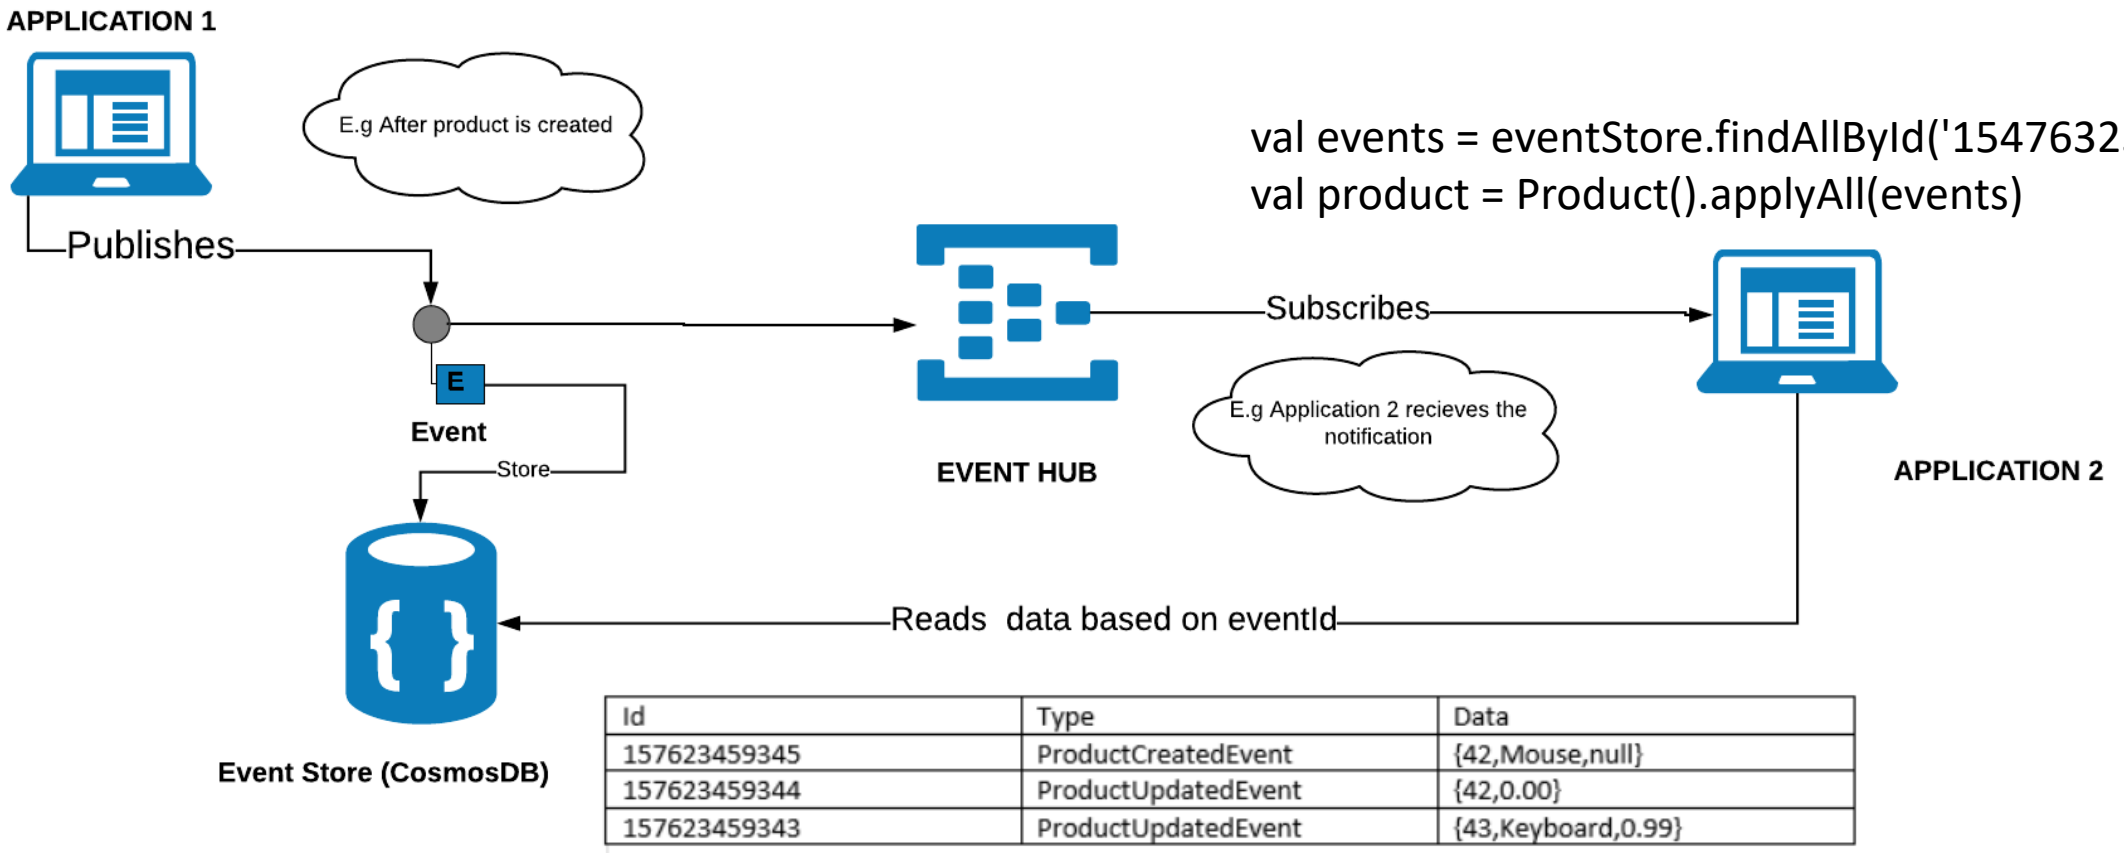


```
val events = eventStore.findAllById('1547632386819');
val product = Product().applyAll(events)
```

Id	Type	Data
157623459345	ProductCreatedEvent	{42,Mouse,null}
157623459344	ProductUpdatedEvent	{42,0.00}
157623459343	ProductUpdatedEvent	{43,Keyboard,0.99}

https://sajeetharan.com/2019/02/03/event-sourcing-with-azure-eventhub-and-cosmosdb/

CQRS

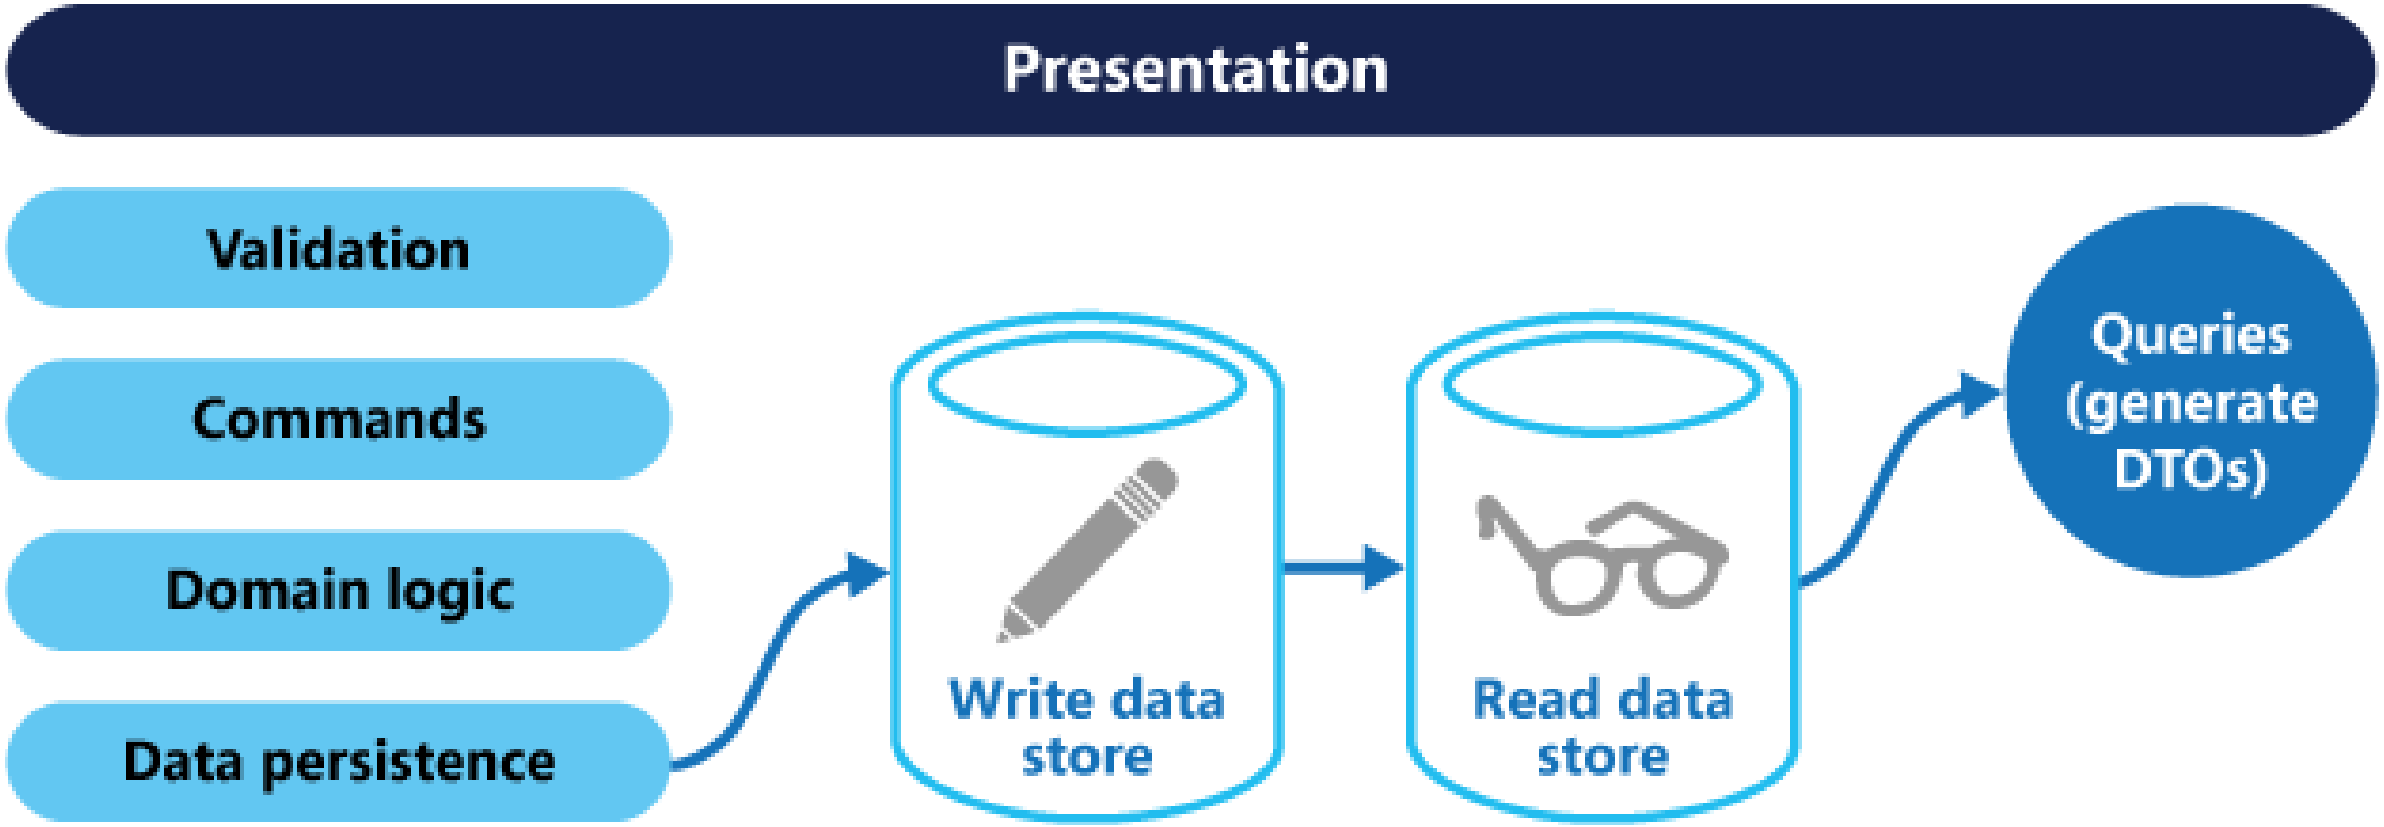
Presentation
Validation
Commands
Domain logic
Data persistence
Write data store
Read data store
Queries (generate DTOs)

What problems does CQRS solve

- CRUD against same entities
- Scaffolding tools optimize for commands
- Columns with diff update frequencies
- Data contention
- Diff Security /Perms

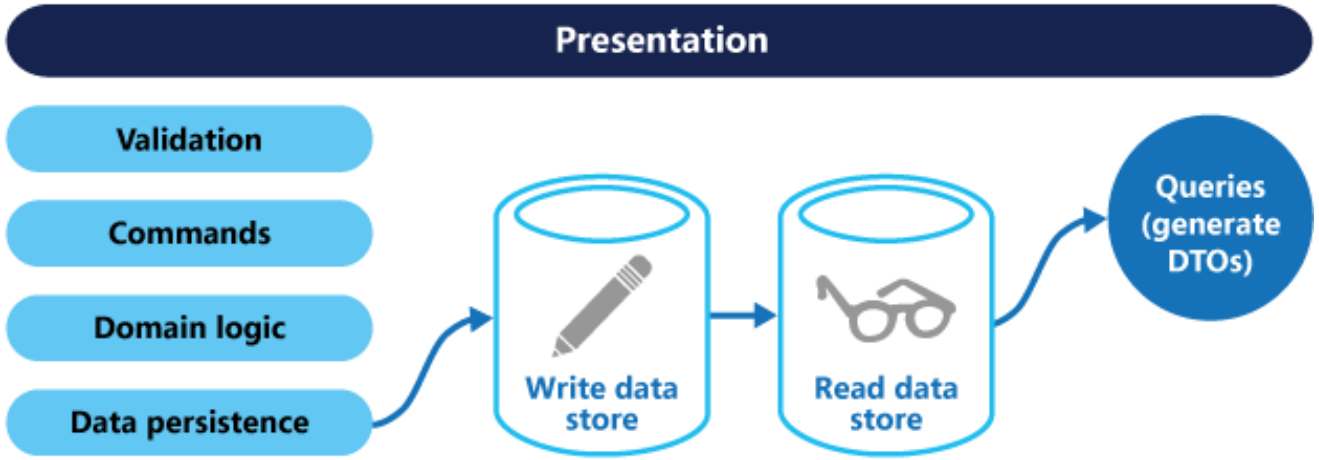

How does CQRS solve these problems?

- Segregate Reads and Commands
- Diff data models
- Separate R/W stores

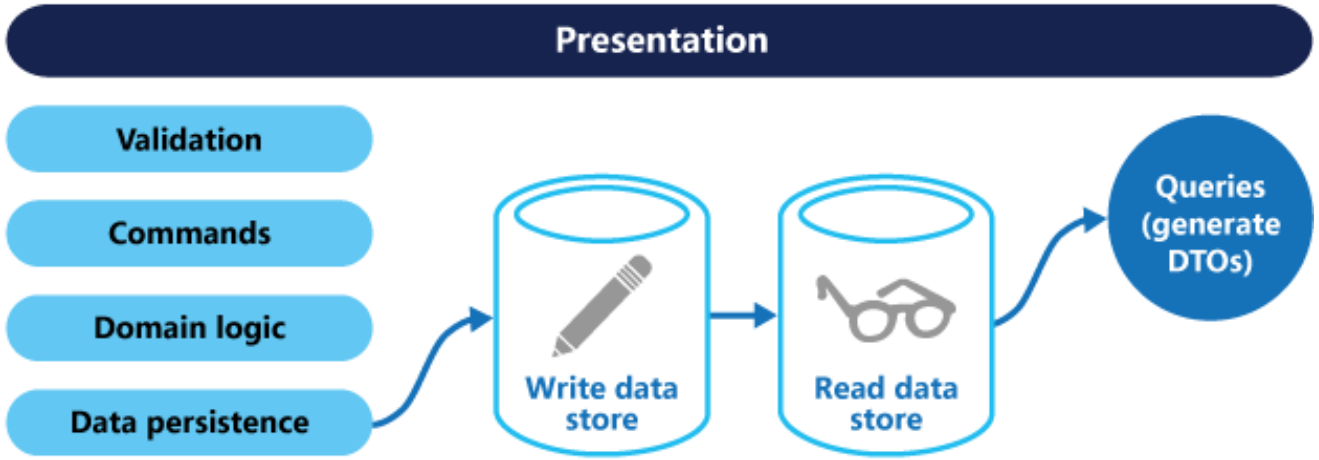

CQRS Considerations

- No Scaffolding
- Increase perf+security
- More complex
- Model change mgmt.
- Limited scope
- Task not data
- **Eventual Consistency**

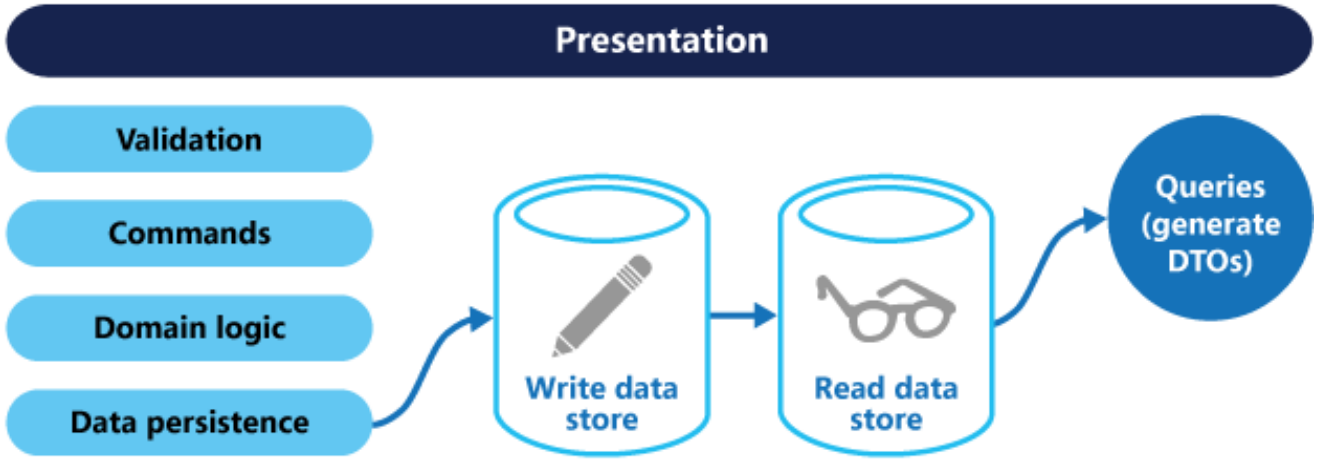

CQRS Use Cases - Good

- Collaborative domain
- DDD
- R to W ratios high
- Separation of concerns
- Different read models
- Evolving BL

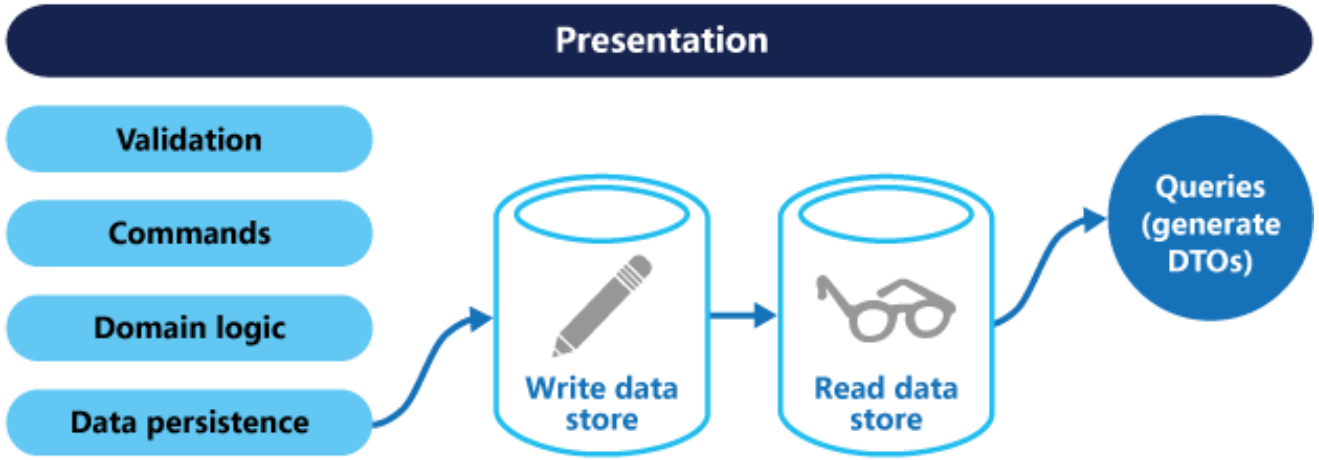

CQRS Use Cases - Bad

- Simple domains
- Less Data
- Whole systems

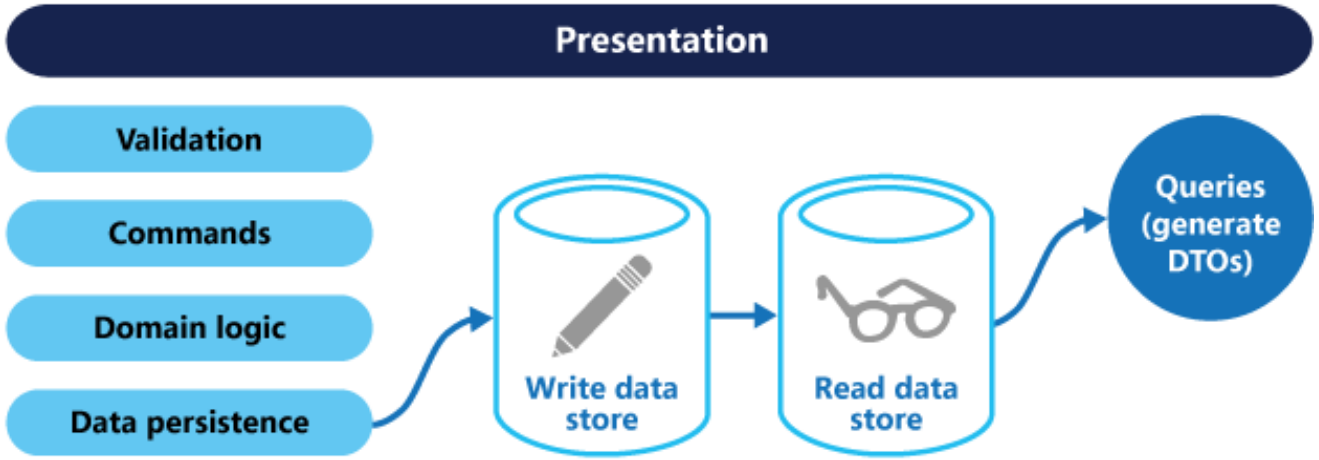

CQRS with Cosmos DB

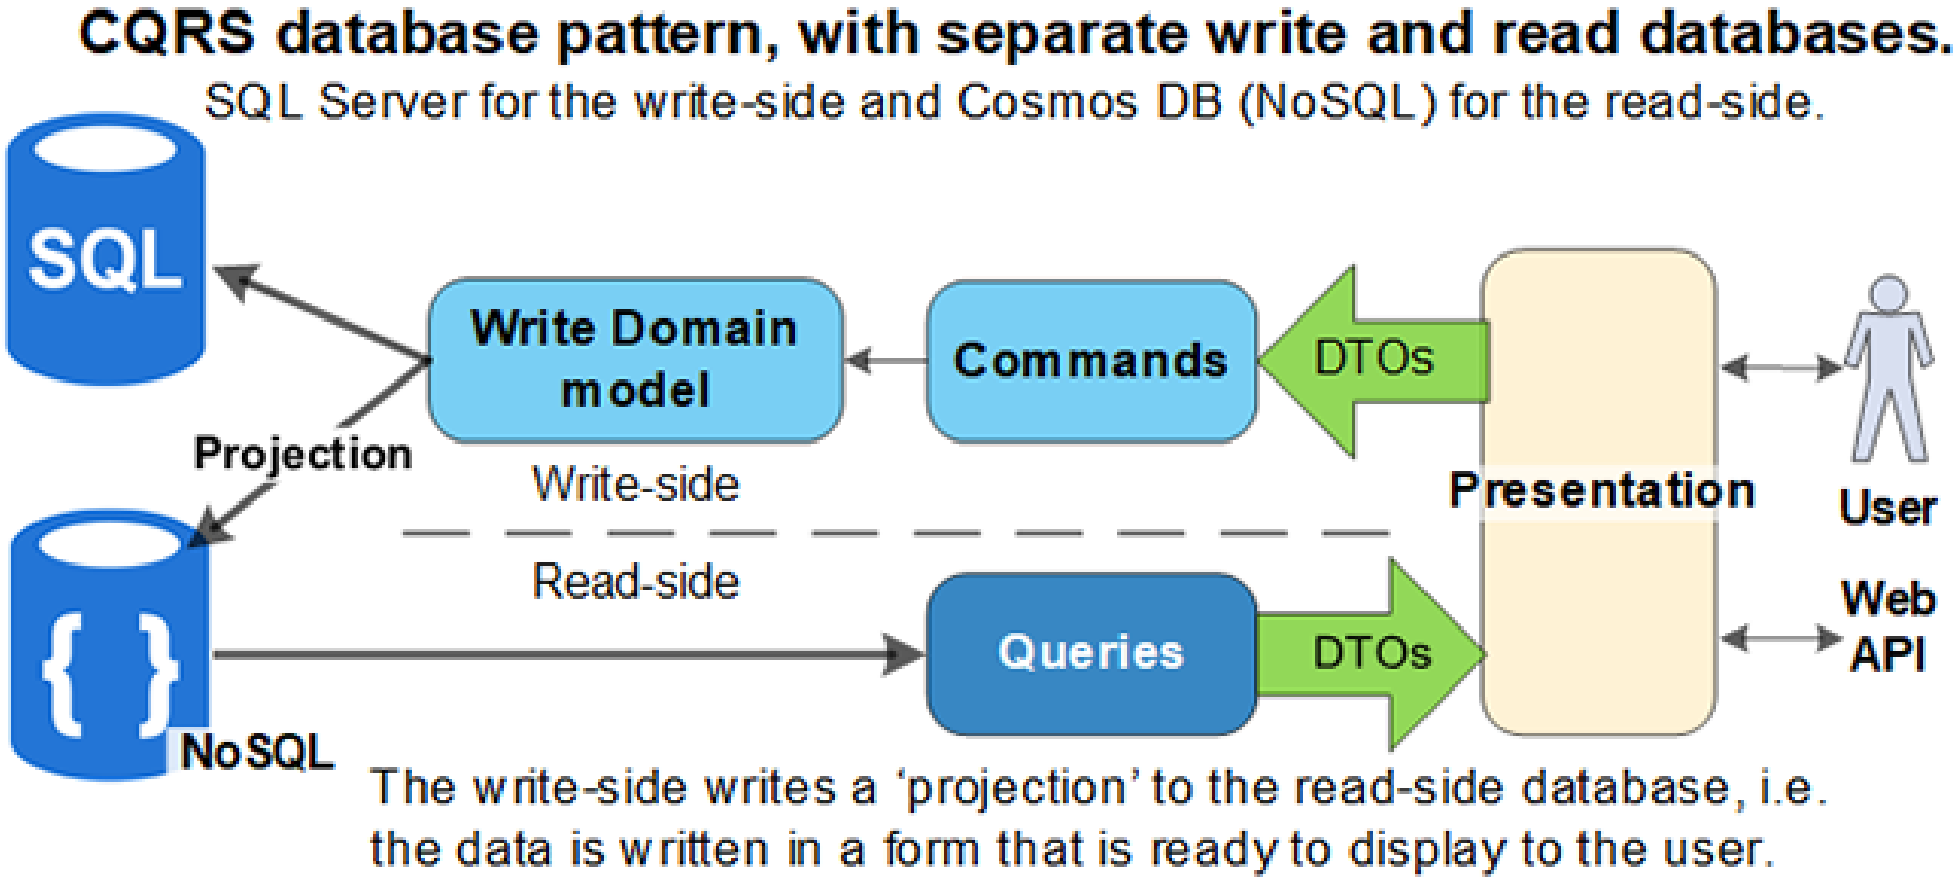

https://www.thereformedprogrammer.net/building-a-robust-cqrs-database-with-ef-core-and-cosmos-db/

How it all ties together

Event sourcing + CQRS + Materialized views

Still Eventually Consistent

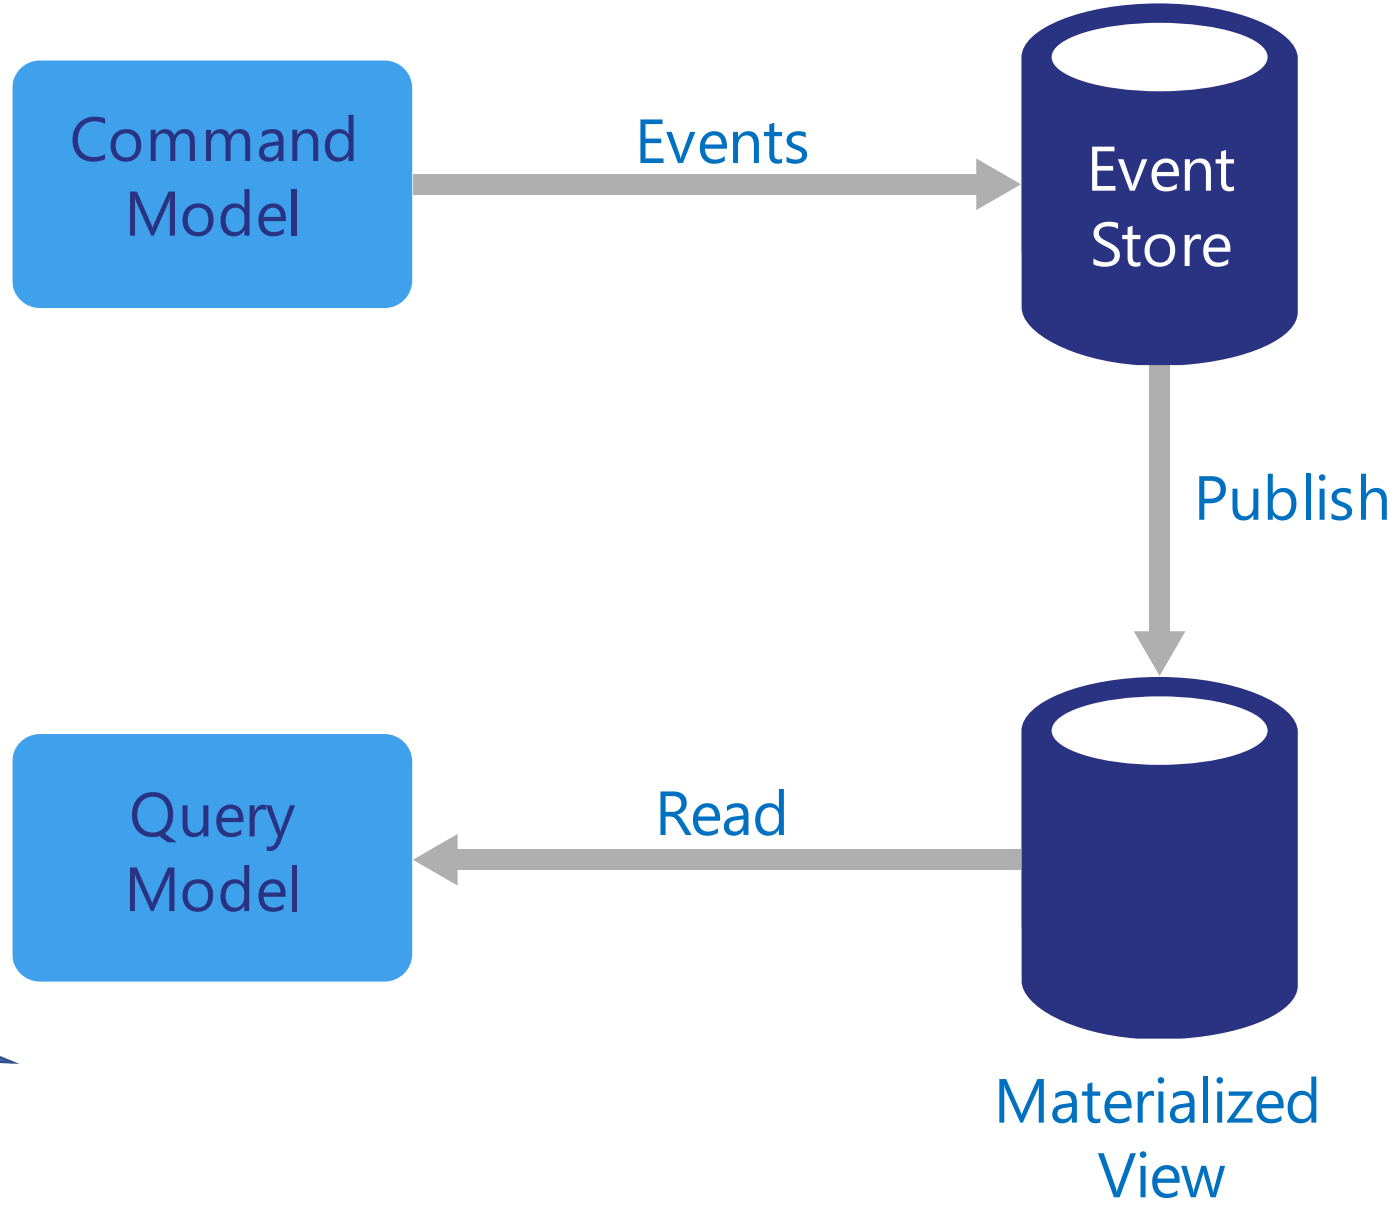

Lambda Architecture

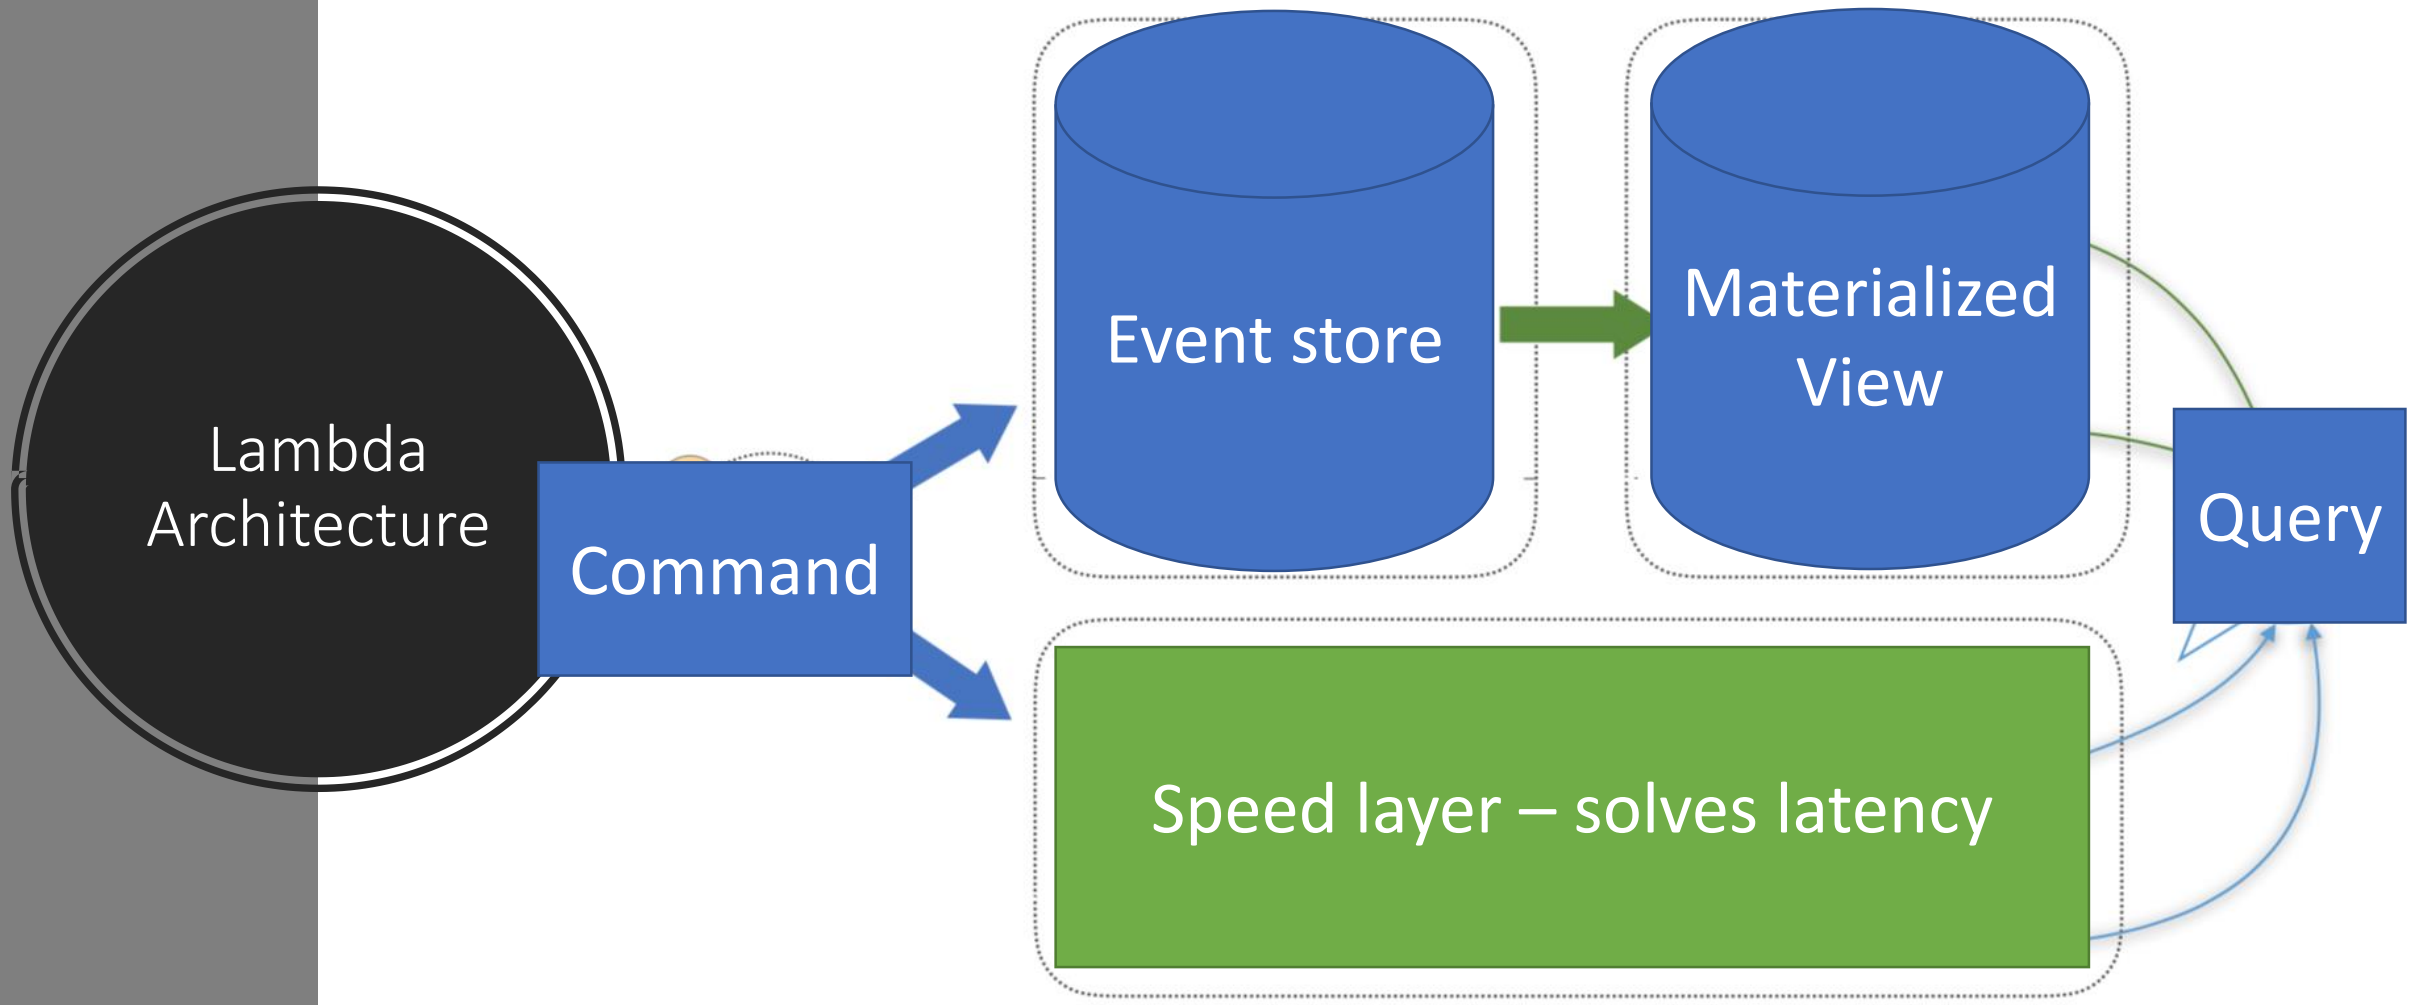

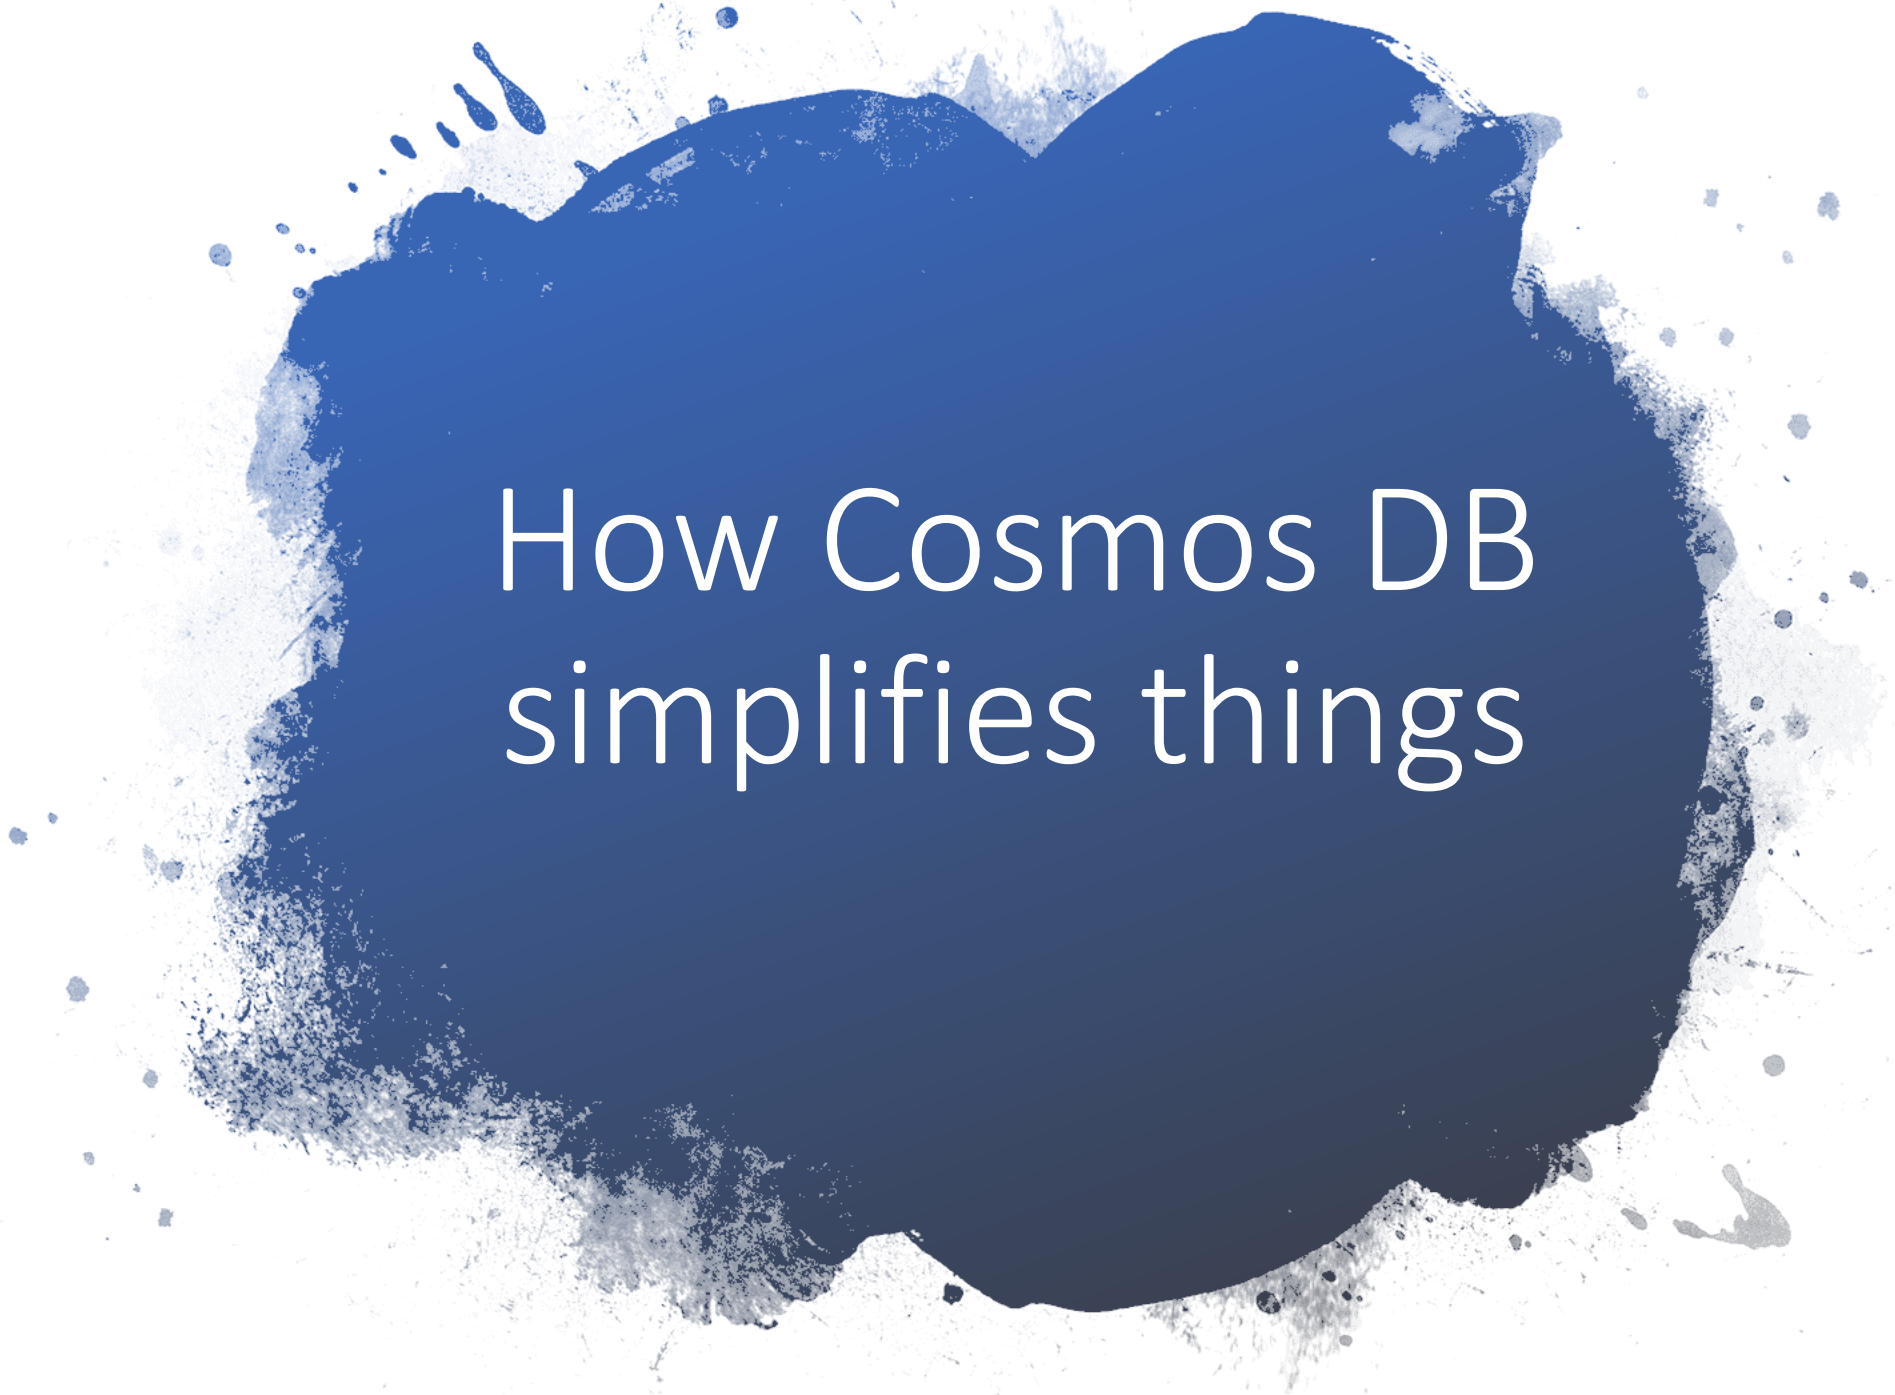
How Cosmos DB
simplifies things

Simplify Lambda Architecture with Cosmos DB

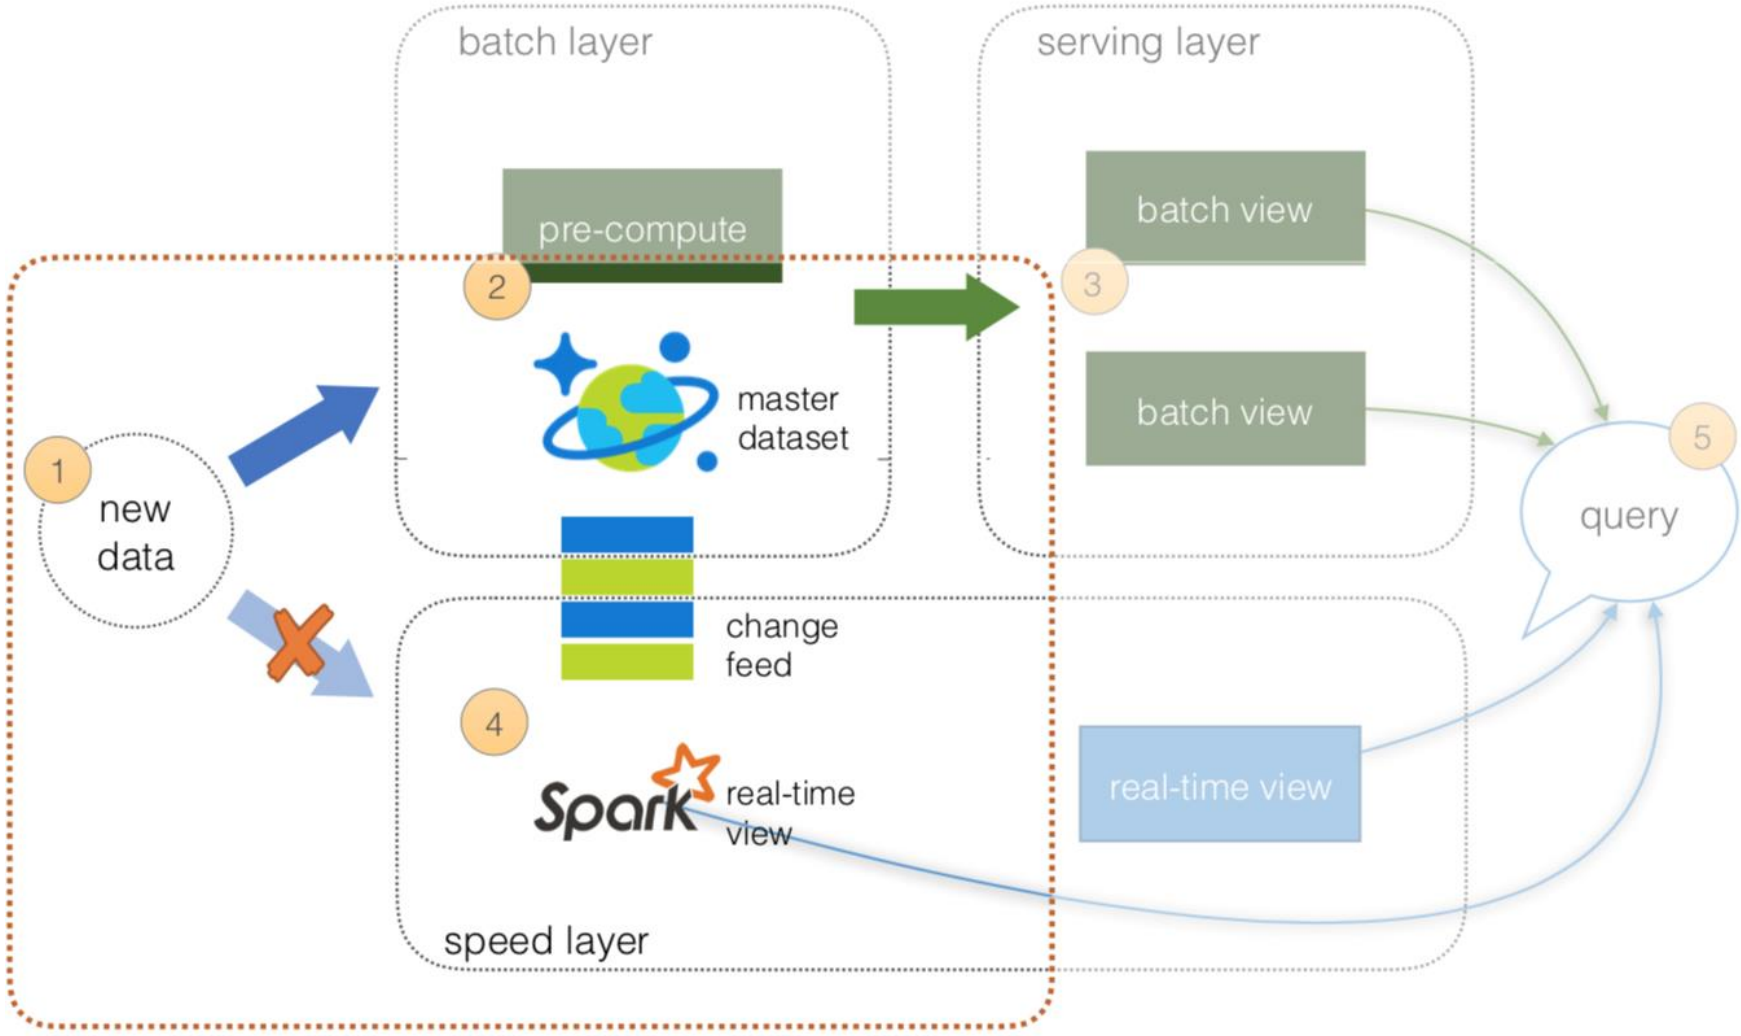

Lambda Architecture Cosmos DB – Batch and Serving Layers

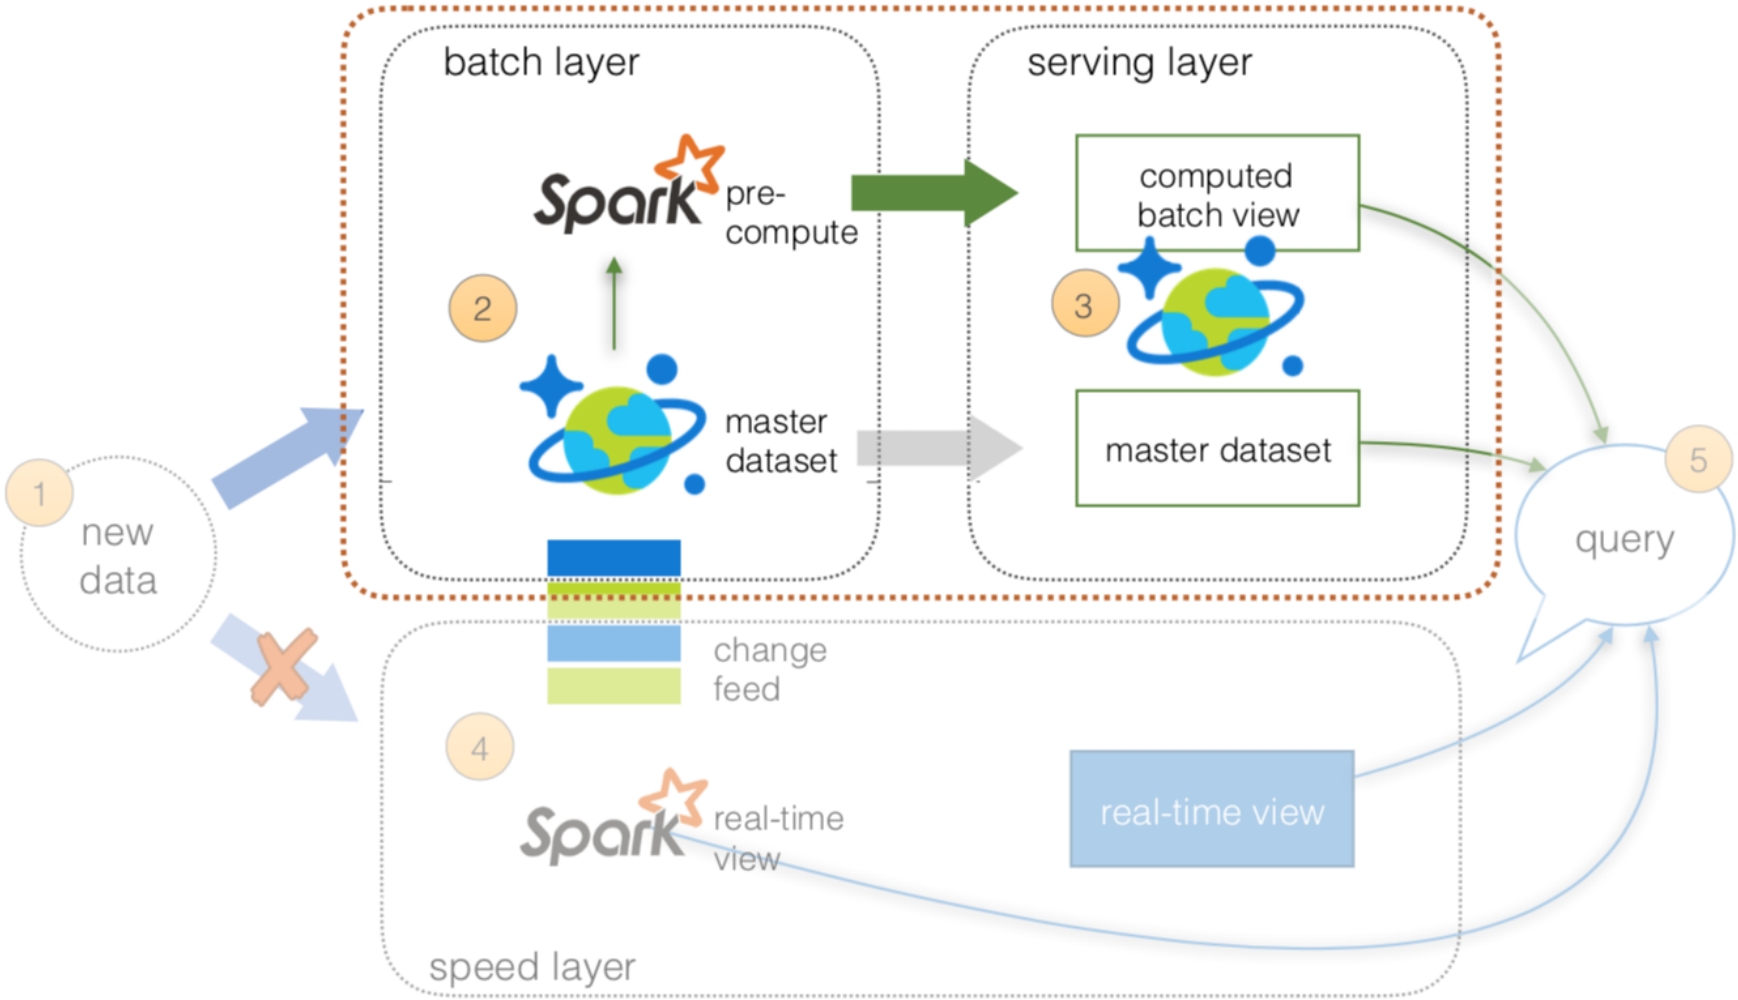

Lambda Architecture Cosmos DB – Speed Layer

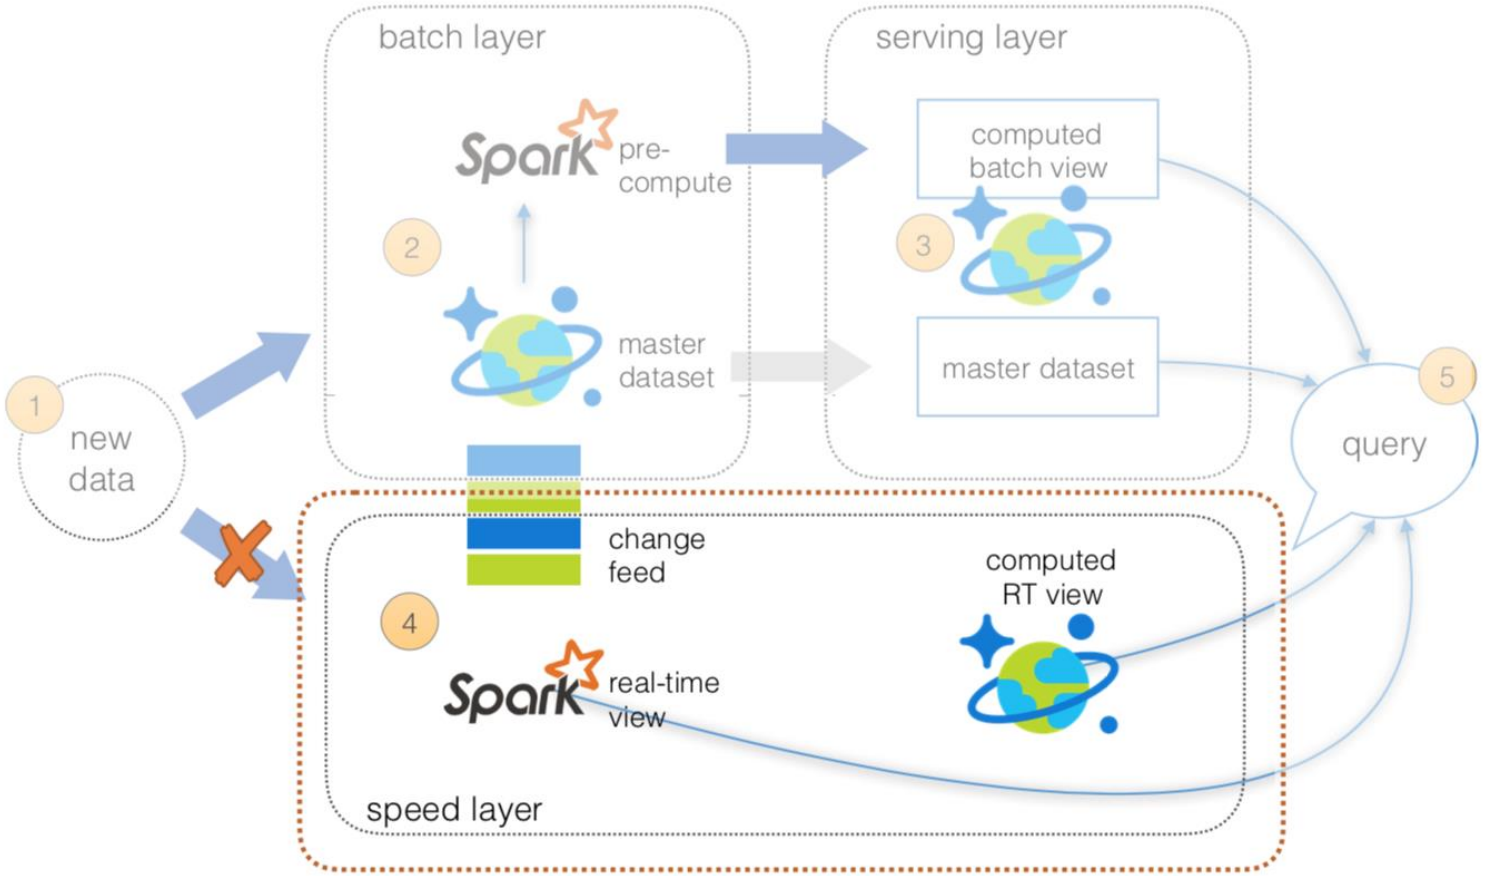

Lambda Architecture Cosmos DB - Re-architected

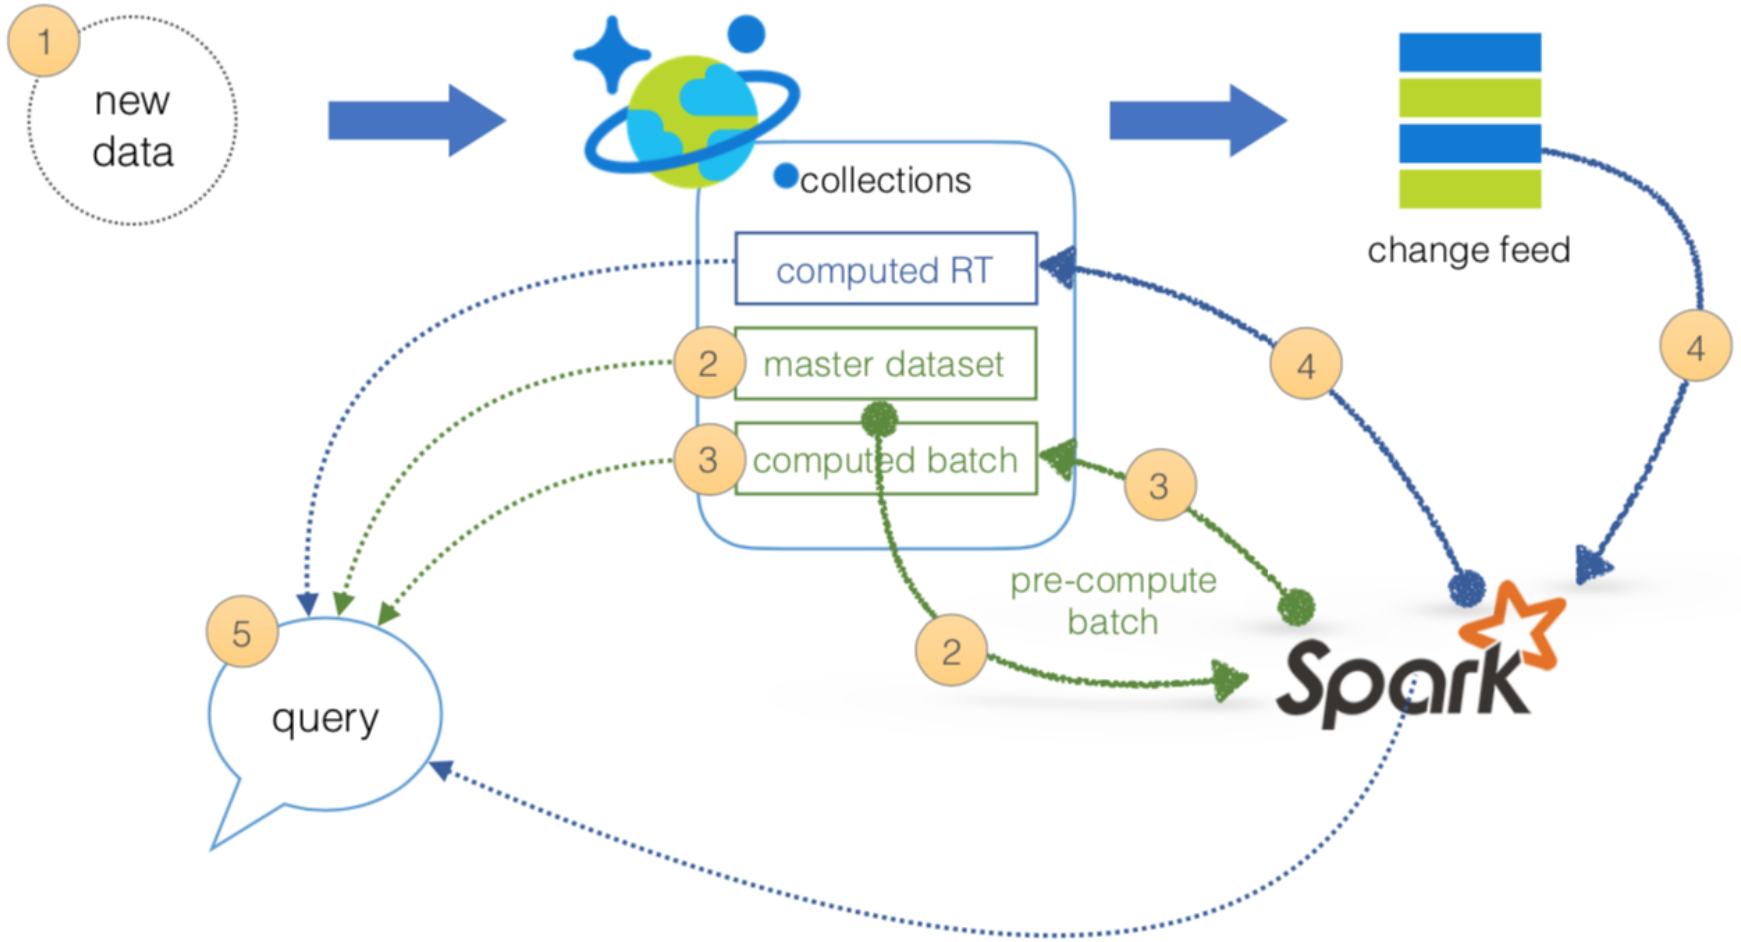

Lambda Architecture - Criticism

Lambda Architecture - downside

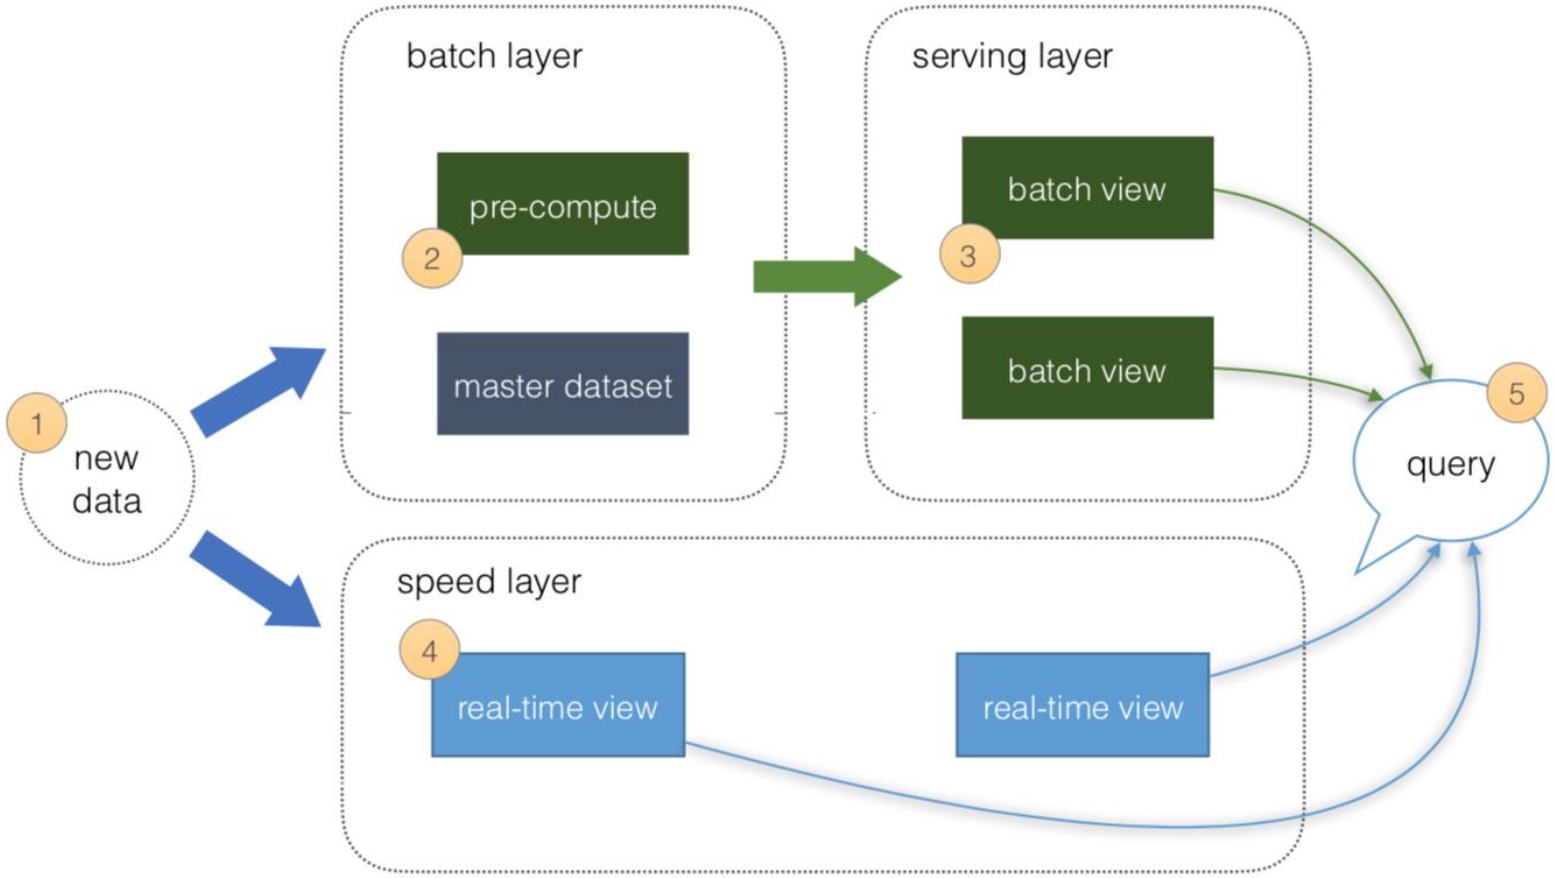

Kappa Architecture

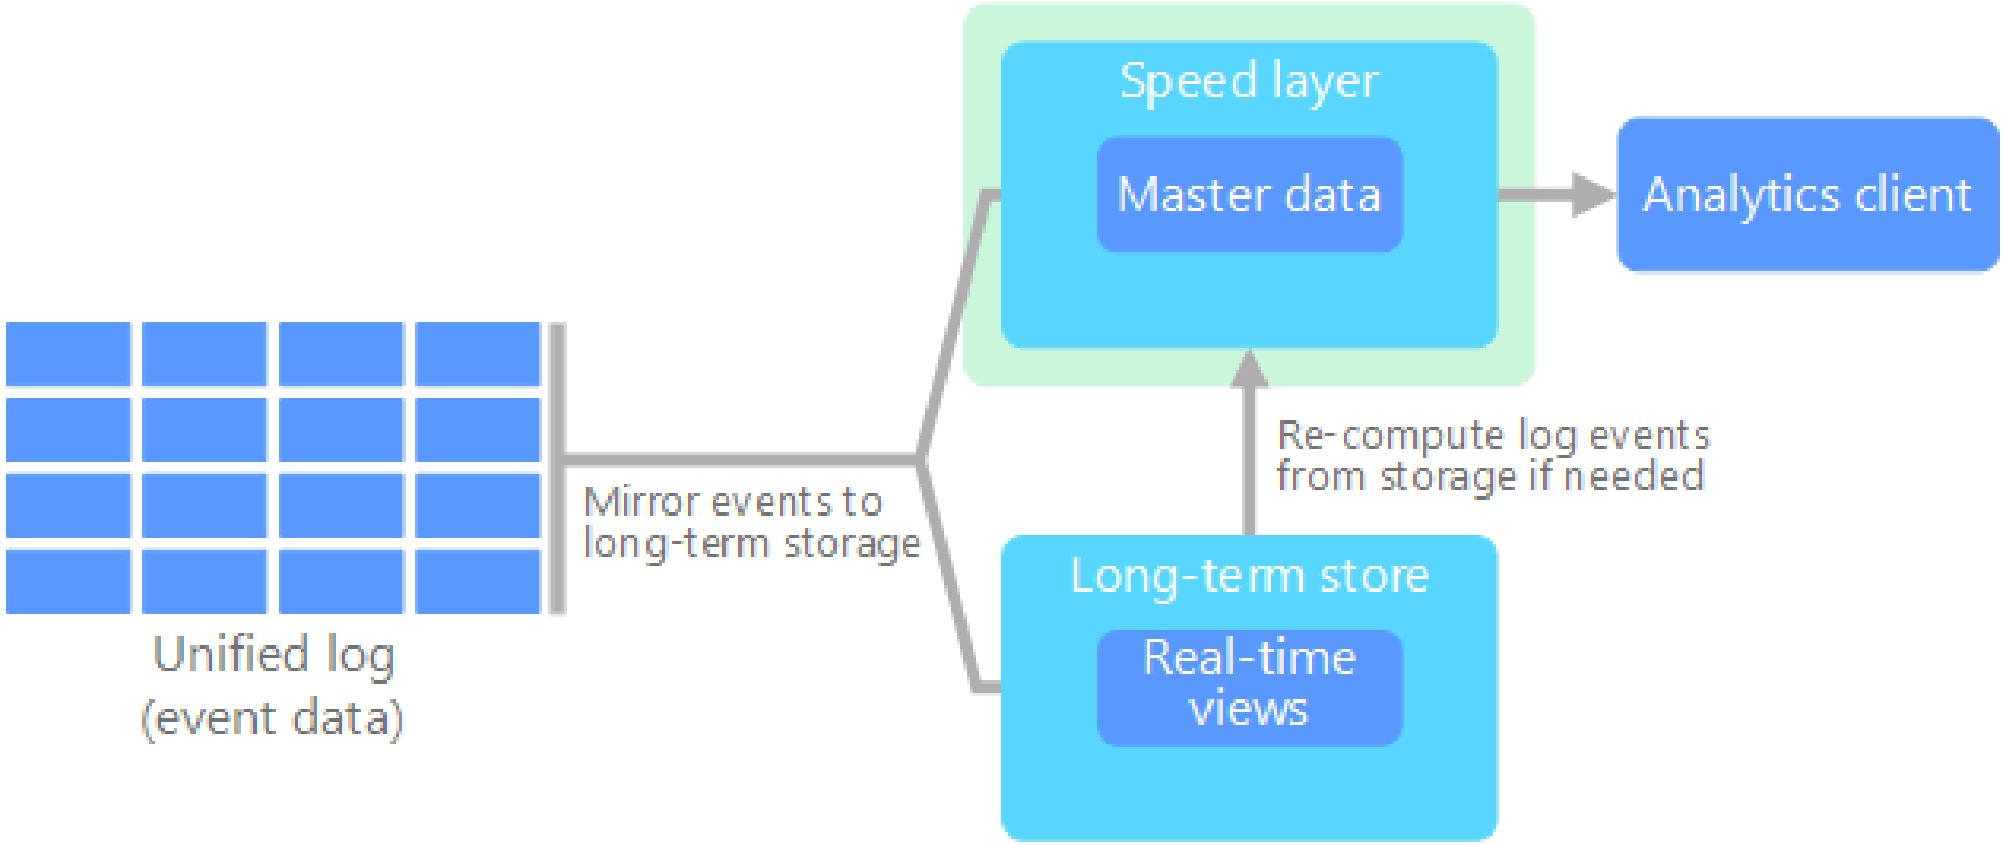

Pros and Cons

CONS

- Re-processing costly
- Speed layer always busy
- Less robust

PROS

- No batch layer, less complex
- Re-processing infrequent
- Runs on fixed memory
- Horizontally scalable
- Fewer resources

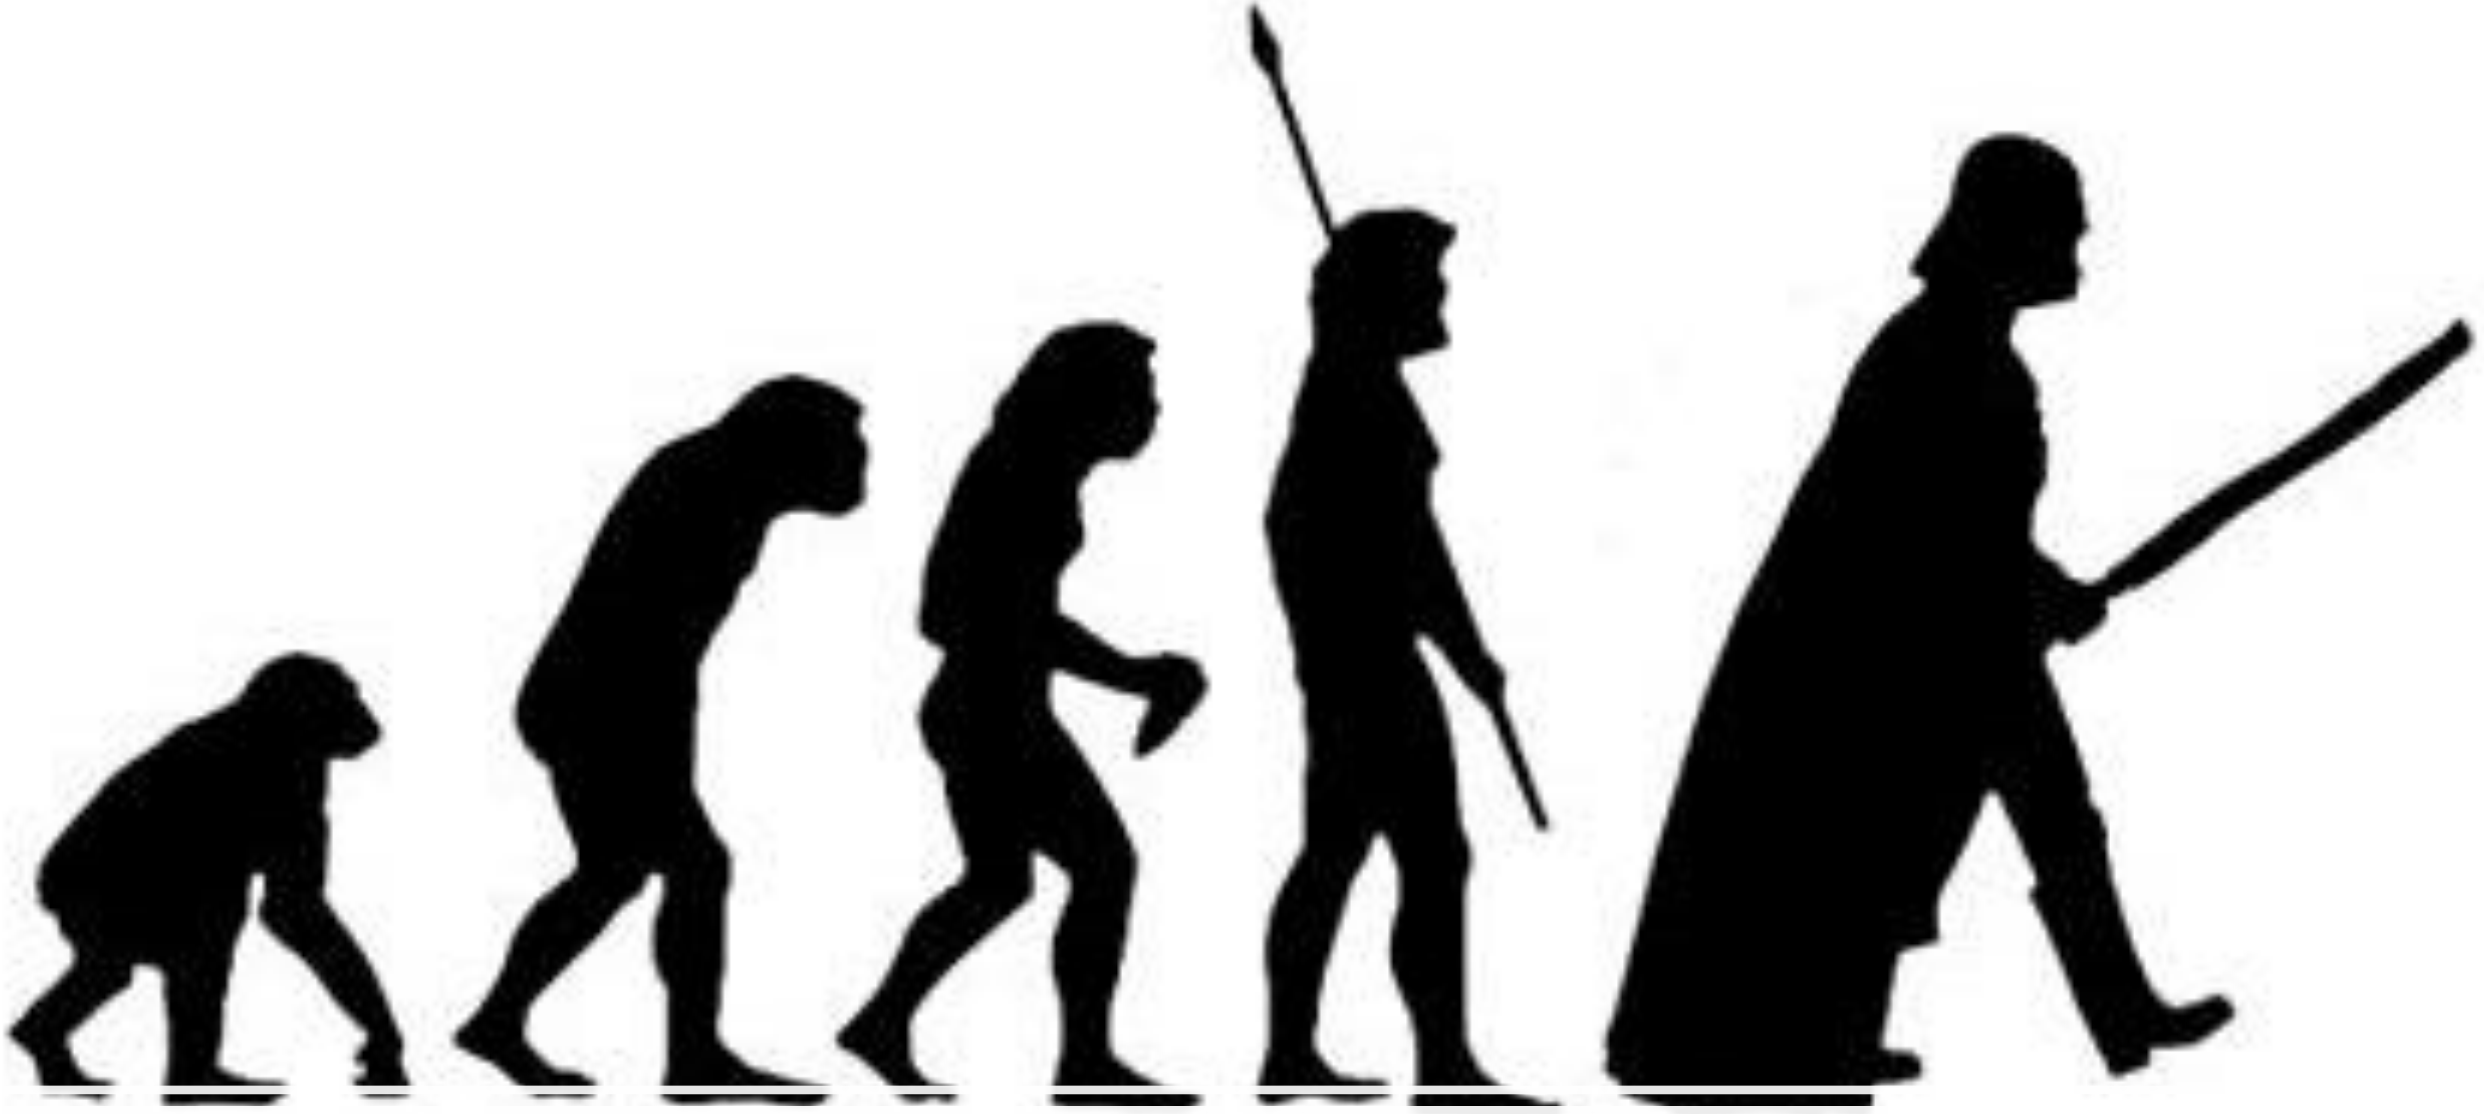

Materialized View

Event Sourcing

CQRS

Kappa/Lambda Architecture

????????

Evolution

Santosh Hari

Azure Consultant @ Nebbia Tech

Azure MVP

President, Orlando .NET UG

Organizer, Orlando Codecamp

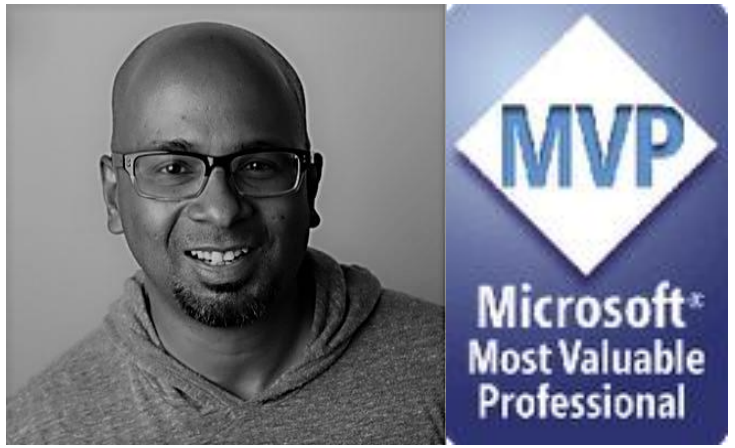

santoshhari.wordpress.com

santosh.hari@newsignature.com

@_s_hari

/in/santoshhari