

SQL is not an excuse
to avoid DevOps

Santosh Hari

Code PaLOUsa 2019 Sponsors

Humana®

XC
Experience
Center



Santosh Hari

Azure Consultant @ Nebbia Tech

Azure MVP

President, Orlando .NET UG

Organizer, Orlando Codecamp



santoshhari.wordpress.com



santosh.hari@newsignature.com



[@_s_hari](https://twitter.com/_s_hari)



[/in/santoshhari](https://in.linkedin.com/in/santoshhari)



Agenda

Problems database professionals experience during deployments

Concerns about DevOps

Why risk avoidance is bad

How DevOps addresses and minimizes risk

Basics of database DevOps

Demo

Parting thoughts

The factory of the future will have only two employees, a **man** and a **dog**. The **man** will be there to feed the **dog**. The **dog** will be there to keep the **man** from touching the equipment.

Warren Bennis - scholar, organizational consultant and author

The title and
some
content
inspired by

acmqueue

[Current Issue](#) [Past Issues](#) [Topics](#)

The May/June 2019 issue of acmqueue is out now

[Subscribers and ACM Professional members login here](#)

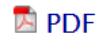


Everything Sysadmin

Development

December 12, 2018

Volume 16, issue 5



PDF

SQL is No Excuse to Avoid DevOps

Automation and a little discipline allow better testing, shorter release cycles, and reduced business risk.

Thomas A. Limoncelli

Yes, I got permission

← **Thomas Limoncelli 6k** ✓
@yesthattom



Hey Thomas, loved your ACM article queue.acm.org/detail.cfm?ref... which is very much relevant to some projects I'm doing currently where I'm literally having to force DevOps with databases included down the throats of my clients. I was planning on doing some conference talks and workshops on this same topic. Would you be ok with me using some portions of your post, especially the title? I'll make sure I attribute. Let me know. Either way, great post. The community needed something like this

Dec 18, 2018, 1:19 PM ✓



Hi! It's an academic publisher. Please feel free to quote and cite appropriately! Best, Tom

Apr 18, 2019, 9:02 PM

Thank you for the response and clarification!!

Apr 19, 2019, 1:07 AM ✓

Current state of database deployments



Deployments take hours if not days



Weeks of negotiation with moving codebases



Infrequent aka big bang



Often performed on weekends



No control over when to deploy your changes

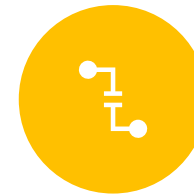
Jumping through hoops for database deployments



Have to perform gymnastics



Lock out users/apps



Take Offline



Disable data writes



Run in parallel and sync up data

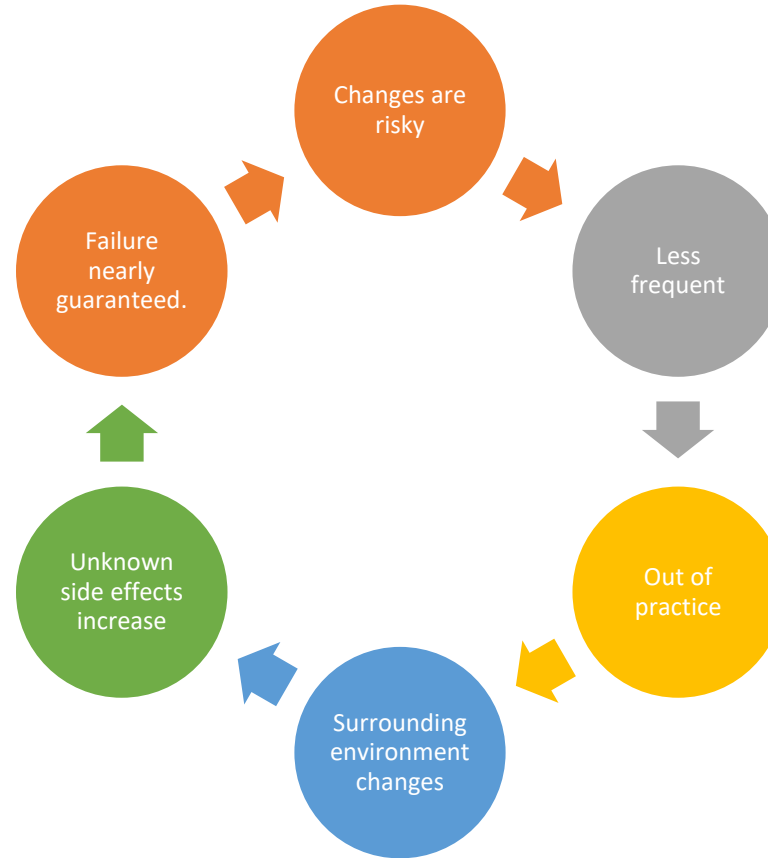
Concerns with DevOps




DEPLOYMENTS TAKE TIME



DEPLOYMENTS ARE RISKY



Risk avoidance causes fragility



DevOps is the union of people,
process, and products to enable
continuous delivery of value to our
end users

Donovan Brown, Cloud Advocate Manager of the
Methods and Practices Org, Microsoft

How to DevOps

Continuous
Integration (CI)

Continuous
Delivery (CD)

Version Control
(VC), mostly Git

Agile planning
and lean project
management

Monitoring and
Logging

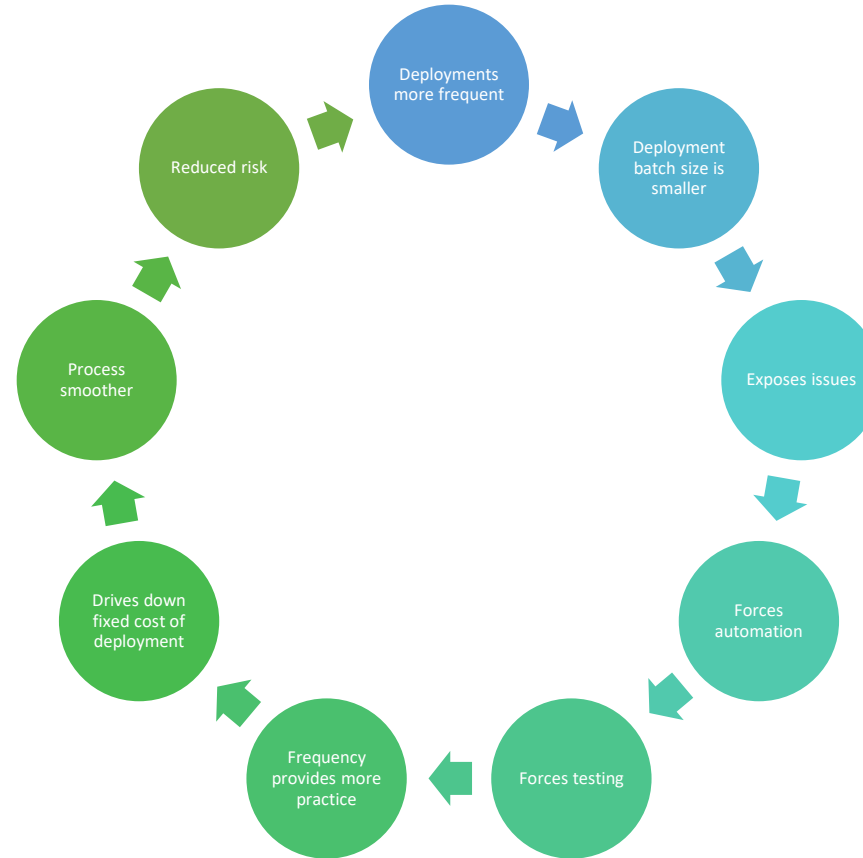
Public and
Hybrid Clouds

Infrastructure-
as-Code (IaC)

Microservices

Containers

Shift Left
philosophy



DevOps is anti-fragile

DevOps for databases

Everything is automated

Schema change management - State-based or migration-based?

Unit tests

Schema validation

Data (Lookup tables, for instance)

Security – firewall, usernames/passwords

Infrastructure and config

Logging and monitoring

Feature Flags

Above all, work closely with the rest of the team

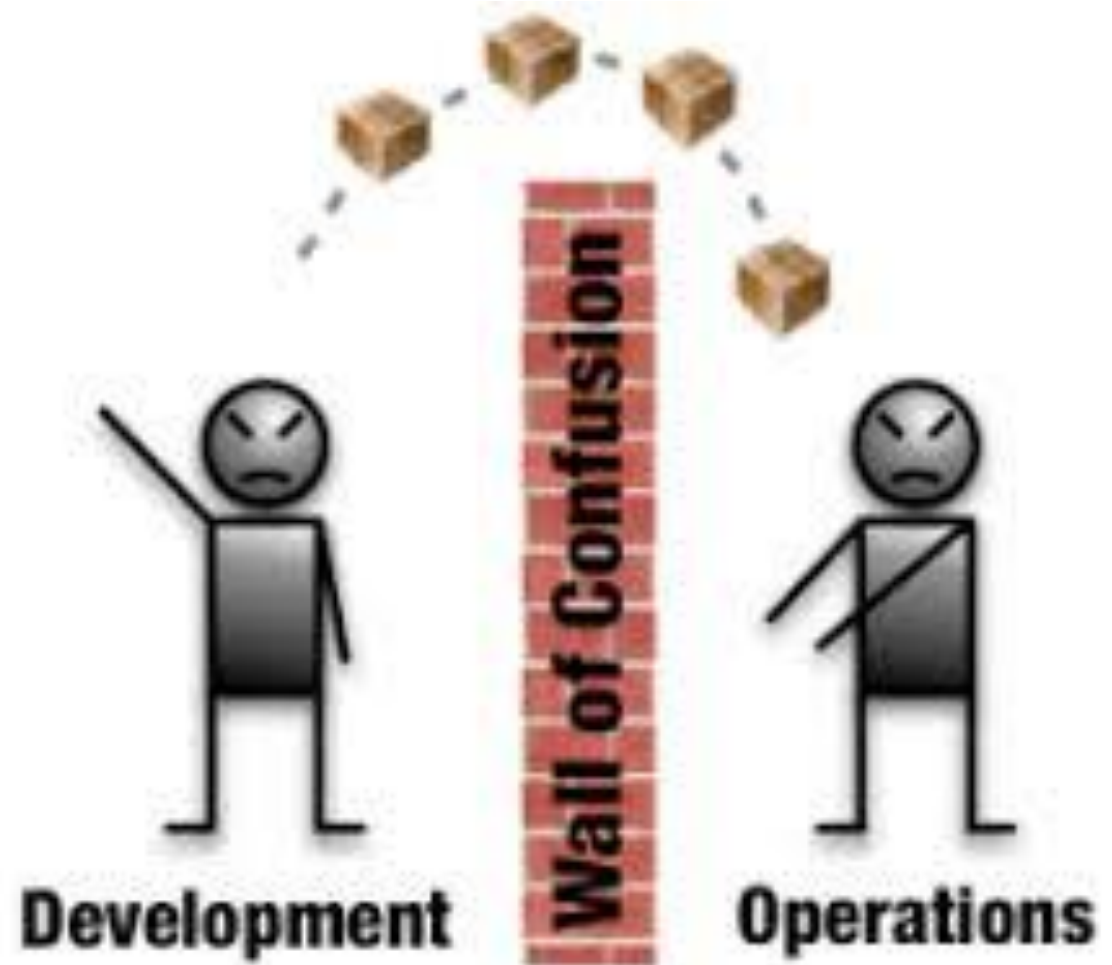
- Deploy database from scratch
- Manage subsequent schema changes
- Infrastructure create/change

Demo

A dark blue, irregular, ink-blot-like shape with splatters on a white background. The shape is roughly circular but has jagged, organic edges. It is surrounded by numerous small, dark blue ink splatters and larger, lighter blue washes that spread outwards from the main shape. The overall effect is that of a fresh ink blot on a clean white surface.

Parting thoughts

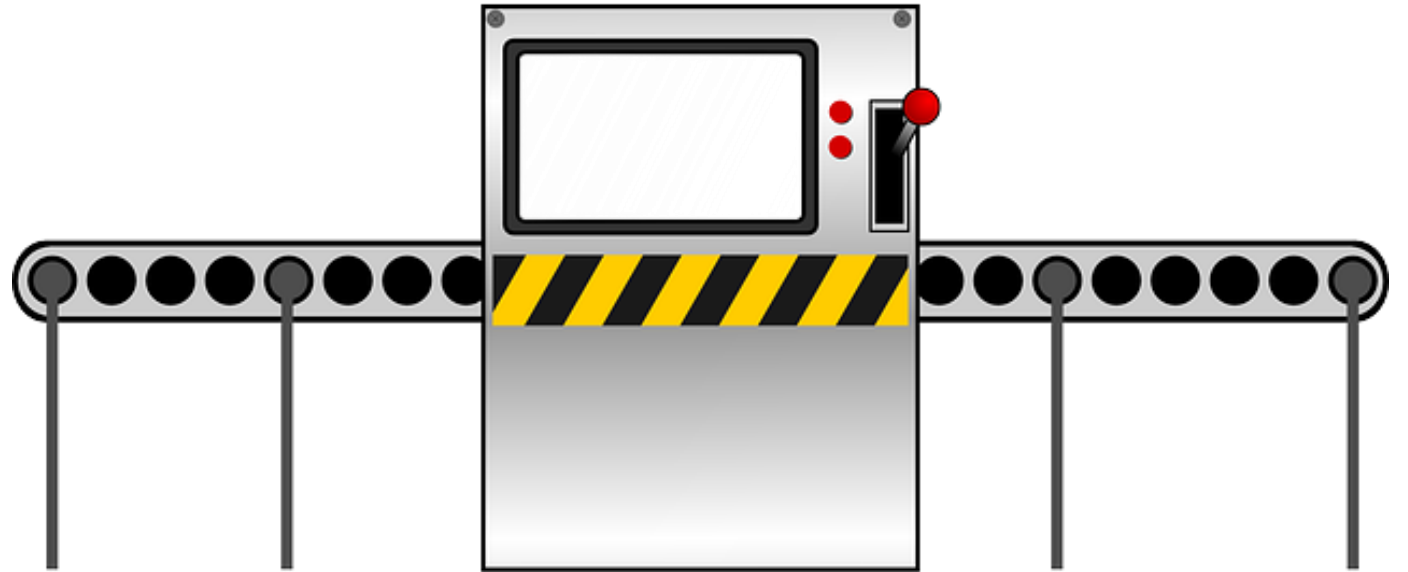
Don't throw
stuff over the
wall



Don't
attempt to
boil the
ocean



Pipeline and
automate
everything



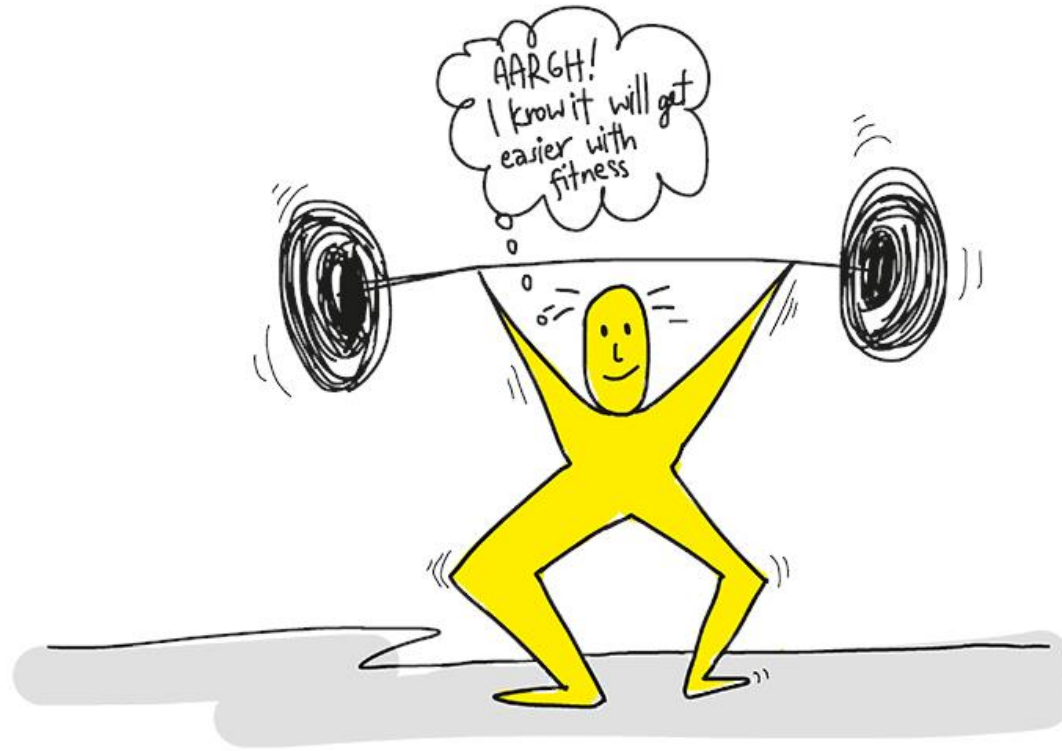
Reduce drift
between
environments



Log and
monitor
everything



DevOps
will hurt at
first



DevOps is a
journey



Resources



SQL is no excuse to avoid DevOps – ACM – Thomas Limoncelli
<https://queue.acm.org/detail.cfm?ref=rss&id=3300018>



Official Microsoft SQL Server DevOps <https://www.microsoft.com/en-us/sql-server/developer-get-started/sql-devops/>



What is DevOps? <https://docs.microsoft.com/en-us/azure/devops/learn/what-is-devops>



Redgate – State based or migration based <https://www.redgate.com/library/state-or-migrations-based-database-development>



The Phoenix Project - <https://www.amazon.com/Phoenix-Project-DevOps-Helping-Business/dp/1942788290>



Accelerate - <https://www.amazon.com/Accelerate-Software-Performing-Technology-Organizations/dp/1942788339>

Santosh Hari

Azure Consultant @ Nebbia Tech

Azure MVP

President, Orlando .NET UG

Organizer, Orlando Codecamp



santoshhari.wordpress.com



santosh.hari@newsignature.com



[@_s_hari](https://twitter.com/_s_hari)



[/in/santoshhari](https://in.linkedin.com/in/santoshhari)



State-based database change management

- SSDT or Redgate SQL Source Control (supports hybrid)
- Single step between current and desired state
- Mostly CREATE statements
- Easy workflow for large teams and complicated databases
- Scripts generated at deployment time using comparison
- Developer involvement minimal
- UI based and automated options available
- Bad for scenarios like column-split or datatype/constraint changes because context is lost

Migration-based database change management

- Migration-first: EF, DbUp, Redgate SQL Change Migration (supports hybrid)
- Captures individual change scripts during dev to effect larger iterative migrations
- Numerically ordered migrations
- Mostly ALTER commands
- Works well with simple data stores for table refactoring or data migration
- Dev and DBA collaborate early in dev cycle to define deployment steps
- Can be executed as code/scripts both manual and through automation
- Last check-in would dominate

Automated Schema Updates

- Automated Schema Updates

Coding for Multiple Schemas

- Coding for Multiple Schemas

Infrastructure and environments

Generate Test Data

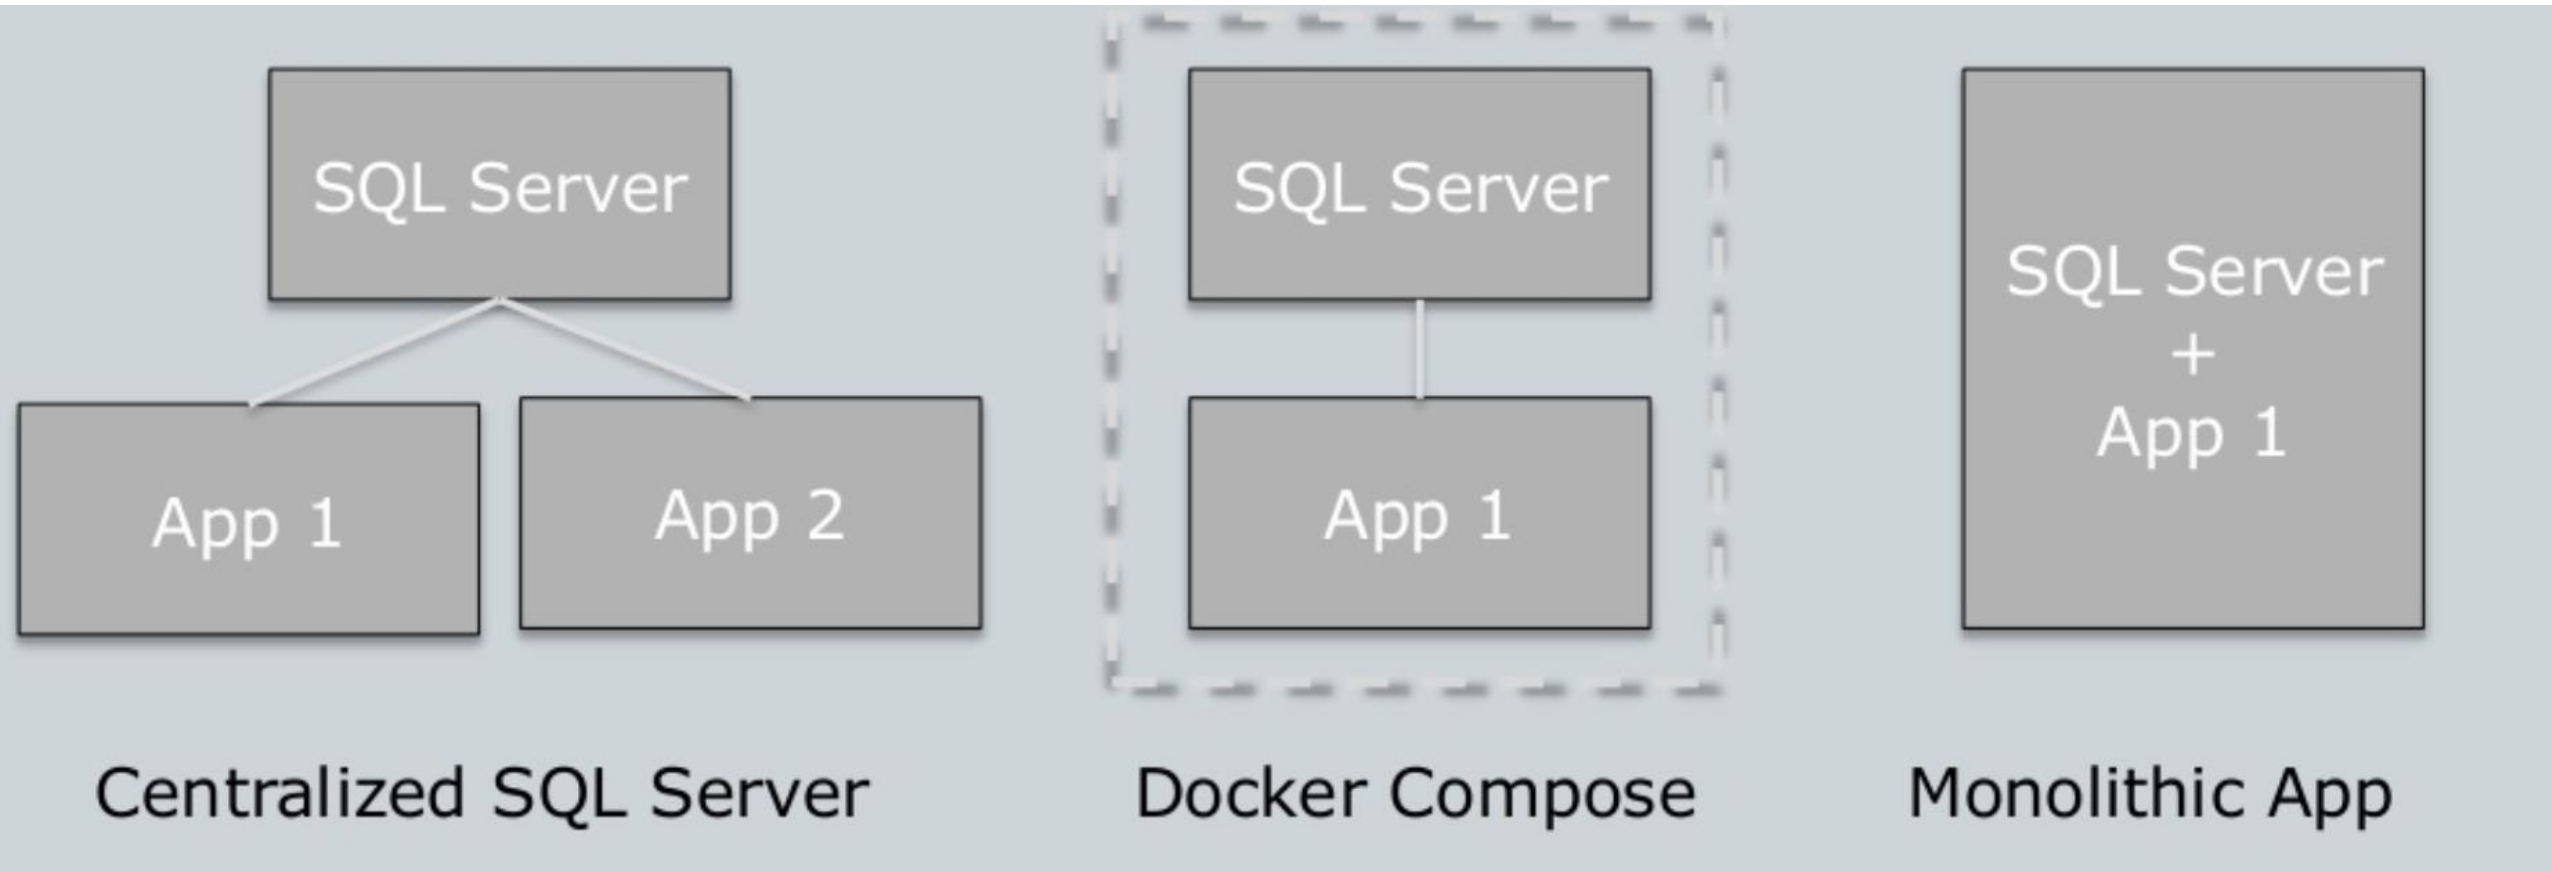
- Generate Test Data

Log and monitor everything

Multi-phase schema commit

- Propose New Schema: checkin new schema, contains both old (primary) and new (backup) fields, no field deletes/renames, DAO abstracts column type changes, make app new schema, write to both old and new columns, migrate old data as needed, easy to rollback
- Flip source of truth: new fields should be primary, old fields will still be valid. Previous release will still work
- Remove old references: Once you're certain you're not going to roll back, stop reading/writing the old data. This is your hard flip.
- requires at least two periods of schema change and one period that involves multiple software rollouts
- requires that your application is well formed. Your application needs to be properly modularized, use reasonable DAOs or otherwise have complete understanding of how the data is accessed

App Deployments with Containers



Santosh Hari

Azure Consultant @ Nebbia Tech


Azure MVP

President, Orlando .NET UG

Organizer, Orlando Codecamp

 santoshhari.wordpress.com

 santosh@nebbiatech.com

 [@_s_hari](https://twitter.com/_s_hari)

 [/in/santoshhari](https://in/santoshhari)

